



Cairo University
Faculty of Engineering
Department of Computer Engineering

Accelera



A Graduation Project Report Submitted
to
Faculty of Engineering, Cairo University
in Partial Fulfillment of the Requirements for the Degree
of
Bachelor of Science in Computer Engineering.

Presented by

Mohamed Ashraf
Mazen Adel

Nesma Osama
Mohamed Khater

Supervised by

Dr. Salma Abdel Monem

April 17, 2026

All rights reserved. This report may not be reproduced in whole or in part, by photocopying or other means, without the permission of the authors or department.

Abstract

Accelerera is a hybrid Python/C++ machine learning framework that combines a high-level Python interface with native execution components. The system is organized around five modules: Accelerera Pipe, the Code Parallelizer, Auto Preprocessing, AutoML, Deployment, and the Benchmark Platform. Accelerera Pipe represents machine-learning workflows as directed graph nodes for preprocessing, modelling, prediction, metric evaluation, branching, and merging. The Code Parallelizer analyzes C/C++ and supported Python code, extracts loop features, predicts parallelization opportunities, and emits OpenMP-annotated C++ code. This report documents the design, methodology, and experimental results available for the implemented modules, while retaining the required report structure for the remaining modules.

Arabic Abstract

Acknowledgment

Table of Contents

Abstract	2
Arabic Abstract	3
Acknowledgment	4
List of Abbreviations	11
List of Symbols	12
1 Introduction	14
1.1 Motivation and Justification	14
1.2 Project Objectives and Problem Definition	14
1.3 Project Outcomes	14
1.4 Document Organization	14
2 Market Visibility Study	15
2.1 Targeted Customers	15
2.2 Market Survey	15
2.2.1 AutoCleanML	15
2.2.2 Kaggle	15
2.3 Business Case and Financial Analysis	16
3 Literature Survey	17
3.1 Background and Previous Work	17
3.1.1 Data Science	17
3.1.2 Machine Learning	17
3.1.3 Supervised Machine Learning	17
3.1.4 Classification	17
3.1.5 Regression	17
3.1.6 Natural Language Processing NLP	17
3.1.7 Exploratory Data Analysis	18
3.1.8 Data Preprocessing	18
3.1.9 Missing Value Imputation	18
3.1.10 Robust Scaling	18
3.1.11 Categorical Encoding	18

3.1.12	Binary Encoding	18
3.1.13	Frequency Encoding	19
3.1.14	Term Frequency-Inverse Document Frequency TF-IDF	19
3.1.15	Data augmentation	19
3.1.16	Classification Evaluation Metrics	19
3.1.16.1	Accuracy	19
3.1.16.2	Precision	20
3.1.16.3	Recall	20
3.1.16.4	F1-score	20
3.1.17	Regression Evaluation Metrics	20
3.1.17.1	Mean Absolute Error	20
3.1.17.2	Mean Squared Error	20
3.1.17.3	R-squared Score	20
3.2	Comparative Study of Previous Work	21
3.2.1	Data Preprocessing and Feature Engineering for Data Mining: Techniques, Tools, and Best Practices [1]	21
3.3	Implemented Approach	21
3.3.1	Auto Preprocessing	21
4	System Design and Architecture	23
4.1	Overview and Assumptions	23
4.1.1	Accelera Pipeline Assumptions	23
4.1.2	Auto Preprocessing Assumptions	23
4.1.2.1	Tabular Dataset Assumptions	23
4.1.2.2	Text Dataset Assumptions	23
4.1.2.3	Image Dataset Assumptions	24
4.1.3	Benchmark Assumptions	24
4.2	System Architecture	25
4.2.1	Block Diagram	25
4.3	Accelera Pipe Module	25
4.3.1	Module Responsibility	25
4.3.2	Graph Execution Model	26
4.3.3	Branching and Node Sharing	27
4.3.4	Memory-Lifetime Optimizations	27
4.3.5	Saving and Loading Pipelines	28
4.4	Code Parallelizer Module	28
4.4.1	Module Responsibility	28
4.4.2	Loop Analysis and Pragma Selection	29
4.4.3	Classifier Model and Rule-Based Decision	30
4.4.4	Python-to-C++ Converter Support	30
4.4.5	Integration with Accelera Pipe	31
4.5	Auto Preprocessing Module	31

4.5.1	Functional Description	35
4.5.1.1	Tabular datasets	35
4.5.1.2	Text datasets	35
4.5.1.3	Image datasets	35
4.5.2	Modular Decomposition	35
4.5.2.1	Tabular Data Preprocessing Functions	35
4.5.2.2	Text Data Preprocessing Functions	38
4.5.2.3	Image Data Preprocessing	39
4.5.2.4	Testing and Reusability:	41
4.5.3	Design Constraints	41
4.5.4	Other Description of the Module	42
4.6	AutoML Module	42
4.7	Deployment Module	42
4.8	Benchmark Platform Module	42
4.8.1	Functional Description	42
4.8.2	Modular Decomposition	43
4.8.2.1	Benchmark Management	43
4.8.2.2	Metric Management	43
4.8.2.3	Submission Management	43
4.8.2.4	Validation and Scoring Scripts	43
4.8.2.5	Leaderboard Interface	44
4.8.3	Design Constraints	44
4.8.4	Other Description of the Module	44
5	System Testing and Verification	45
5.1	Testing Setup	45
5.2	Testing Plan and Strategy	45
5.2.1	Module Testing	45
5.3	Accelera Pipe Module	45
5.4	Code Parallelizer Module	47
5.4.0.1	Auto Preprocessing Module	50
5.4.0.2	AutoML Module	52
5.4.0.3	Deployment Module	52
5.4.0.4	Benchmark Platform Module	52
5.4.1	Integration Testing	52
5.4.2	Testing Schedule	52
5.4.3	Comparative Results to Previous Work	52
6	Conclusions and Future Work	53
6.1	Faced Challenges	53
6.2	Gained Experience	53
6.3	Conclusions	53

6.4 Future Work 53

References 54

A Development Platforms and Tools 56

B Use Cases 57

C User Guide 58

D Code Documentation 59

E Feasibility Study 60

List of Figures

- 4.1 Before duplicate first-node sharing 27
- 4.2 After duplicate first-node sharing 27
- 4.3 Code Parallelizer module workflow 29
- 4.4 Integration flow for compiled preprocessing inside Accelera Pipe 31
- 4.5 Tabular Preprocessing Workflow 32
- 4.6 Text Preprocessing Workflow 33
- 4.7 Image Preprocessing Workflow 34

List of Tables

4.1	Accelera Pipe builder methods	26
4.2	Supported operations in the Python-to-C++ converter	30
4.3	Detected Column Types in Tabular Data	37
5.1	Accelera Pipe latency scaling in seconds	46
5.2	Accelera Pipe RSS memory scaling in MB	46
5.3	Selected candidate and accuracy across sample sizes	47
5.4	Largest Accelera Pipe comparison	47
5.5	OpenMP classifier report	48
5.6	Hard parallelizer evaluation	48
5.7	Linux parallelizer and Accelera Pipe integration benchmark	48
5.8	Linux direct example-script verification runs	49
5.9	Windows direct example correctness and path timing	49
5.10	Windows compilation and process overhead	50
5.11	Datasets Used in Module Testing	51
5.12	Models and Evaluation Metrics Used in Module Testing	52

List of Abbreviations

AI	Artificial Intelligence
AST	Abstract Syntax Tree
CPU	Central Processing Unit
ML	Machine Learning
OpenMP	Open Multi-Processing
RSS	Resident Set Size
SVC	Support Vector Classifier
VMS	Virtual Memory Size

List of Symbols

X	Input feature matrix
y	Target vector
P	Source program
L	Extracted loop
$F(L)$	Feature representation of loop L

Contacts

Team Members

Name	Email	Phone Number
Name of Student 1	abc1@email.com	+2 01xxxxxxxxx
Nesma Osama	nesma.ahmed03@eng-st.cu.edu.eg	+2 01067193062
Name of Student 3	abc1@email.com	+2 01xxxxxxxxx
Name of Student 4	abc1@email.com	+2 01xxxxxxxxx

Supervisor

Name	Email	Number
Salma Abdelmoniem	abc1@email.com	+2 01114302362

Chapter 1: Introduction

1.1 Motivation and Justification

Accelera addresses the practical gap between the productivity of Python and the performance requirements of machine-learning workflows and loop-heavy data processing. A sequential script evaluates alternatives one after another and leaves memory lifetime and parallel execution decisions to the user. The project provides a graph-based pipeline backend and a code parallelization path to make these execution decisions explicit and reusable.

1.2 Project Objectives and Problem Definition

The project objective is to provide an integrated framework for constructing, executing, evaluating, and accelerating data-science workflows. The problem is to preserve a convenient Python-level programming model while enabling native graph scheduling, controlled intermediate-data lifetime, and OpenMP-backed execution for supported custom code.

1.3 Project Outcomes

The project is organized into the following modules:

- Accelera Pipe for graph-based machine-learning workflows.
- Code Parallelizer for source analysis and OpenMP generation.
- Auto Preprocessing for tabular, text, and image data preparation.
- AutoML for model selection and optimization.
- Deployment and Benchmark Platform components for serving and evaluation.

1.4 Document Organization

Chapter 2 reserves the required market visibility study. Chapter 3 contains the literature-survey structure. Chapter 4 documents the system design and implemented module methodology. Chapter 5 presents testing and verification results. Chapter 6 concludes the report and records future work. The appendices retain the required development, user-guide, and feasibility-study sections.

Chapter 2: Market Visibility Study

In this chapter, we'll identify Accelera's target customers, talk about similar tools and platforms in the market, discuss their pros and cons, and provide the business case study and financial analysis of the system.

2.1 Targeted Customers

Our targeted customers are data scientists and machine learning engineers, both beginners and experts. Accelera aims to support these users by providing a tool that automates and accelerates several steps in the data science life cycle. This helps users reduce repetitive manual work, and save time.

2.2 Market Survey

We found that many existing platforms provide useful features for data science and machine learning development, but these features are often available in separate tools or are limited to specific parts of the workflow. However, Accelera aims to provide these features within one toolkit.

2.2.1 AutoCleanML

It is a Python tool mainly focused on automating data cleaning and preprocessing tasks for machine learning datasets.

Pros:

- Makes machine learning data preparation automatic and smart.
- Provides a report that explains the preprocessing decisions made by the tool.

Cons:

- Only designed for cleaning and preparing data..
- It does not support preprocessing for image and text datasets.

2.2.2 Kaggle

Kaggle is an online platform where the data science and machine learning community comes together. It is best known for its competition platform.

Pros:

- It introduces competitions.

- It ranks users in competitions.
- It has many public datasets.

Cons:

- It mainly depends on the notebooks style, so projects can get hard to organize.
- It doesn't provide a complete machine learning workflow in a single toolkit.

2.3 Business Case and Financial Analysis

Chapter 3: Literature Survey

In this chapter, we give an overview of the topics, algorithms, and techniques needed to understand the project as a whole. We discuss the technical subjects and background knowledge, and then summarize the research papers related to our project.

3.1 Background and Previous Work

This section explains the concepts and techniques used in Accelerera units.

3.1.1 Data Science

Data science is a field that extracts knowledge and insights from data using scientific methods, processes, and algorithms.

3.1.2 Machine Learning

Machine learning is a field of artificial intelligence that lets computers learn patterns from data without needing to be explicitly programmed. A machine learning model is trained using training data , and then it uses what it learned to make predictions on new data.

3.1.3 Supervised Machine Learning

The training dataset includes a label column that represents the target output. This column is used during training to learn the mapping between the input features and the expected result.

3.1.4 Classification

Classification is a supervised machine learning task whose output is a label or category.

3.1.5 Regression

Regression is a supervised machine learning task whose output is a continuous numerical value.

3.1.6 Natural Language Processing NLP

Natural language processing is a field that focuses on enabling computers to understand and work with human language. Text data can't be used directly by most machine learning models, so it needs to be cleaned and turned into numerical features

first.

3.1.7 Exploratory Data Analysis

Exploratory Data Analysis is an important step in data analysis where we explore, summarize, and visualize data to understand its structure, detect patterns, and check relationships between variables.

3.1.8 Data Preprocessing

Data preprocessing is the stage of transforming raw data into a clean and usable format for machine learning models.

3.1.9 Missing Value Imputation

Missing value imputation is a preprocessing technique used when some values in a dataset are empty. Instead of removing all rows that contain missing values, the missing values can be replaced using a suitable value.

3.1.10 Robust Scaling

Robust scaling is a scaling technique that uses the median and the interquartile range instead of the mean and standard deviation.

$$x' = \frac{x - \text{median}(X)}{\text{IQR}(X)}$$

where x' is the scaled value, x is the original value, $\text{median}(X)$ is the median of the feature values, and $\text{IQR}(X)$ is the interquartile range.

3.1.11 Categorical Encoding

Categorical encoding is the process of converting text categories into numeric values. Many traditional machine learning models and preprocessing pipelines expect numerical input.

3.1.12 Binary Encoding

Binary encoding is a categorical encoding technique that first gives each category a number, then converts this number into binary form. The binary digits are then stored as separate columns.

3.1.13 Frequency Encoding

Frequency encoding replaces each category with how often it appears in the dataset. If a category appears many times, it receives a higher frequency value.

$$f(c) = \frac{\text{count}(c)}{N}$$

where $f(c)$ is the frequency of category c , $\text{count}(c)$ is the number of rows containing this category, and N is the total number of rows.

3.1.14 Term Frequency-Inverse Document Frequency TF-IDF

It is a technique used to convert text into numerical features by giving each word a weight based on how important it is in a document and across the whole dataset. Term Frequency measures how often a word appears in a document. Inverse Document Frequency gives lower weight to words that appear in many documents, because these words are usually less useful for classification.

$$TFIDF(t, d) = TF(t, d) \times IDF(t)$$

where $TFIDF(t, d)$ is the weight of term t in document d , $TF(t, d)$ is the frequency of the term in the document, and $IDF(t)$ measures how rare the term is across all documents.

3.1.15 Data augmentation

Is a technique used to increase diversity of a dataset without actually collecting new data. It works by applying various transformations to the existing data to create new, modified versions of data that helps the model generalize better.

3.1.16 Classification Evaluation Metrics

Classification metrics are used to evaluate the performance of a classification model.

3.1.16.1 Accuracy

Accuracy measures how many predictions are correct out of all predictions:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP is true positives, TN is true negatives, FP is false positives, and FN is false negatives.

3.1.16.2 Precision

Precision measures how many predicted positive samples are actually positive:

$$Precision = \frac{TP}{TP + FP}$$

3.1.16.3 Recall

Recall measures how many actual positive samples were correctly detected:

$$Recall = \frac{TP}{TP + FN}$$

3.1.16.4 F1-score

The F1-score combines precision and recall:

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

3.1.17 Regression Evaluation Metrics

Regression metrics are used when the model predicts continuous numerical values.

3.1.17.1 Mean Absolute Error

Mean Absolute Error measures the average absolute difference between the real values and predicted values:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

where y_i is the actual value and $\hat{y}_i$ is the predicted value.

3.1.17.2 Mean Squared Error

Mean Squared Error gives more penalty to large errors:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

3.1.17.3 R-squared Score

R-squared score measures how well the regression model explains the variation in the actual values.

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Where y_i is the actual value, $\hat{y}_i$ is the predicted value, and $\bar{y}$ is the mean of the actual values.

3.2 Comparative Study of Previous Work

3.2.1 Data Preprocessing and Feature Engineering for Data Mining: Techniques, Tools, and Best Practices [1]

This paper shows the role of data preprocessing and feature engineering in machine learning and data mining projects. It says that preprocessing is not only a preparation step before model training, but also an important part of the complete machine learning pipeline. It also presents a general preprocessing framework that includes data cleaning, data transformation, feature construction, feature selection, and dimensionality reduction. For each stage, it explains different techniques and discusses the advantages and disadvantages of each one.

3.3 Implemented Approach

3.3.1 Auto Preprocessing

In this module we decided to implement an automated preprocessing approach for tabular datasets. The approach was designed to prepare data for machine learning models by applying suitable preprocessing steps automatically according to the type and characteristics of each column and some user configurations.

The general framework of the approach is as follows:

1. Take the dataset and some user configurations.
2. Remove duplicate records.
3. Split the data into training and validation sets.
4. For each column, we get this information: data type, percentage of missing values, unique values, and basic stats.
5. Drop unnecessary columns, such as constant columns or columns with missing values above the selected threshold.
6. Detect column types automatically.
7. Make Exploratory Data Analysis.
8. Apply missing value imputation.
9. Apply categorical encoding.
10. Apply robust scaling.
11. Apply feature selection using mutual information.

12. Automatically saves all needed files for inference.
13. Generate a report.

The chosen approach follows a classical automated preprocessing pipeline. The main preprocessing decisions are:

- **Missing values:** Median imputation is used for numerical features because it is simple and more robust to outliers than mean imputation. For categorical features, most-frequent imputation is used because it replaces missing values with valid existing categories.
- **Numerical scaling:** `RobustScaler` is used for numerical features because tabular datasets may contain outliers or skewed values. It depends on the median and interquartile range, which makes it more stable than standard scaling when extreme values exist.
- **Categorical encoding:** The encoding method is selected according to the feature type and cardinality. Binary columns are encoded using ordinal encoding, low-cardinality categorical columns are encoded using binary encoding, and high-cardinality categorical columns are encoded using frequency encoding. This helps reduce the number of generated features and avoids high dimensionality.
- **Feature selection:** Mutual information is used because it is a filter-based method that can work with both classification and regression tasks. It helps keep the most informative features and remove weak features before model training.

This approach was chosen because it gives a good balance between automation, simplicity, and speed. More advanced techniques, such as K-nearest neighbour imputation, autoencoders, and wrapper feature selection, can be useful in specific cases, but they need more computation.

One of the important additions in this project is that the module is not only used for data preparation. It also supports exploratory data analysis report.

Also the module was extended to support text and image data preparation, not only tabular data.

Chapter 4: System Design and Architecture

4.1 Overview and Assumptions

This section introduces the assumptions made for each module in Accelera and explains why these assumptions were needed during the implementation.

4.1.1 Accelera Pipeline Assumptions

The system combines Python-facing APIs with native execution components. The design assumes that machine-learning workflows can be represented as a graph of operations and that only a restricted, analyzable subset of custom source code should be translated to parallel C++.

4.1.2 Auto Preprocessing Assumptions

For the Auto Preprocessing module, the assumptions depend on the type of dataset being processed. The module supports three types of data: tabular, text, and image datasets.

4.1.2.1 Tabular Dataset Assumptions

For tabular datasets, the Auto Preprocessing module assumes that:

- The dataset consists of numerical and categorical features, and the task is either classification or regression. Because we try to implement the standard structured machine learning problems where supervised learning is applied..
- The user provides configuration parameters to control preprocessing operations. This allows the pipeline to adapt to different datasets and user requirements.

4.1.2.2 Text Dataset Assumptions

For text datasets, the Auto Preprocessing module assumes that:

- we assume that the dataset has one text column and one categorical target label. This makes the pipeline easier to apply and keeps the number of features smaller after converting the text into numerical values.
- The user provides configuration parameters for text cleaning and feature extraction. To give the user a level of control over the preprocessing pipeline.

4.1.2.3 Image Dataset Assumptions

- The dataset is for image classification and follows this folder structure.

```

Training/
  Cats/
  Dogs/

Validation/
  Cats/
  Dogs/
  
```

- All images are stored in supported formats and are accessible. This ensures compatibility with image loading and preprocessing libraries.
- The user provides configuration parameters for preprocessing and augmentation. Because different datasets require different preprocessing strategies such as resizing, normalization, and augmentation.

4.1.3 Benchmark Assumptions

For the Benchmark module, the following assumptions are made. These assumptions are introduced due to system design constraints, particularly the limited server storage capacity for handling large dataset uploads.

- When the user creates a benchmark, they provide a tabular dataset. This ensures that the module operates on structured data suitable for standard machine learning evaluation.
- All datasets are stored in publicly accessible Google Drive links. This design choice is adopted to reduce server storage requirements.
- The user must provide three CSV files:
 - A training dataset containing both features and target labels.
 - A testing dataset containing the same feature structure as the training set but without the target column.
 - A ground truth target file containing only the target labels for the test set.

This separation ensures a standard evaluation protocol and prevents data leakage during benchmarking.

- When submitting predictions, the user provides a Google Drive link to a CSV file containing the predicted results. This ensures consistency in submission format across all users.

- The submitted prediction file must have the same length and format as the expected output to ensure correct evaluation and fair comparison between different users' models.
- When creating a benchmark, it is assumed that the user selects an evaluation metric that is appropriate for the problem type and provides a correct metric configuration. This is necessary to guarantee meaningful and valid performance measurement.
- For admin users, when defining an evaluation metric, it is assumed that a valid metric from the Scikit-learn library is selected. Because we use Scikit-learn metrics for evaluation.
- The selected metric must be properly configured and is expected to return a single scalar value for performance comparison and ranking purposes.
- The user is allowed to have only one active submission per benchmark. The user cannot create multiple submissions, but they are allowed to update their existing submission. This constraint is enforced to reduce storage usage.

4.2 System Architecture

4.2.1 Block Diagram

The implemented system is organized around graph execution, source-code parallelization, automated preprocessing, model-selection, deployment, and benchmarking components. The following module sections describe the material available in the previous thesis.

4.3 Accelera Pipe Module

4.3.1 Module Responsibility

Accelera Pipe is implemented under `accelera/src/accelera_pipe/`. Its responsibility is to provide a Python builder API over a C++ graph backend. The user writes a pipeline in Python, but the graph structure, node scheduling, branch execution, and intermediate memory lifetime are handled by the native backend.

The core files are:

- `core/pipeline_base.py`: owns or receives a native `graph.Graph` instance and provides shared save/load behavior.
- `core/pipeline.py`: implements the public builder API: `preprocess`, `model`, `predict`, `metric`, `merge`, and `branch`.

- `core/executed_graph.py`: wraps a trained graph for repeated inference.
- `core/metric.py`: defines the metric abstraction.
- `wrappers/`: contains supervised and unsupervised metric adapters.

4.3.2 Graph Execution Model

Each pipeline operation becomes a graph node. A preprocessing node transforms the input matrix, a model node fits or applies an estimator, a prediction node performs inference, a metric node evaluates the prediction, and a merge node combines multiple branch outputs. The graph is built incrementally, so the Python API remains fluent while the backend receives enough structure to execute independent branches in parallel.

Table 4.1: Accelera Pipe builder methods

Method	Node type	Main role
<code>preprocess(name, func)</code>	PREPROCESS	Transform input data
<code>model(name, model)</code>	MODEL	Fit or apply estimator
<code>predict(name, test_data)</code>	PREDICT	Run model prediction
<code>metric(name, metric)</code>	METRIC	Score predictions
<code>merge(name, method)</code>	MERGE	Combine branch outputs
<code>branch(name, ...)</code>	Multiple	Create alternative graph paths

The pipeline returns two objects after training: the metric or prediction results and an ExecutedGraph. The executed graph contains the fitted objects required for inference. This separation is important because training uses temporary arrays, validation targets, and candidate branches, while inference should keep only the selected executable state.

Algorithm 4.1: Pipeline training and executed graph construction

Input: training data X , target y , graph G , selection strategy S

Output: results R , executed graph E

- 1: Attach X and y to the input node of G
- 2: Execute ready preprocessing, model, prediction, metric, and merge nodes
- 3: For each candidate path:
- 4: collect validation metric and path metadata
- 5: Select path P according to strategy S
- 6: Clone the selected fitted nodes from P into a compact executed graph E
- 7: Remove temporary training arrays and non-required execution data from G
- 8: Return results R and executed graph E

4.3.3 Branching and Node Sharing

Branching is one of the main reasons Accelera Pipe uses a graph rather than a linear pipeline object. A user may compare several preprocessors, several models, or complete candidate sub-pipelines. The backend can then schedule independent branches concurrently. A key optimization added to the branch construction step is duplicate first-node sharing. If a branch starts with the same single preprocessing object as another branch, the graph creates the node once and connects the later branch consumers to the same output. This avoids running identical first-stage transformations multiple times.

The sharing rule is intentionally conservative. It is applied only to single branch nodes, the first node of a branch list, or a list containing one item. Later duplicated nodes inside two different lists are not merged automatically, because by that point their inputs may already be different. For example, $[p_1, p_2, p_3]$ and $[p_4, p_2, p_3]$ do not share p_2 ; the first path feeds p_2 with $p_1(X)$, while the second feeds it with $p_4(X)$.




Figure 4.1: Before duplicate first-node sharing




Figure 4.2: After duplicate first-node sharing

4.3.4 Memory-Lifetime Optimizations

The largest memory pressure in Accelera Pipe appears during branch execution. Each preprocessing node can produce a transformed training array, and that array may be consumed by several downstream model nodes. The backend therefore tracks how many consumers still need each intermediate output. When the last consumer finishes, the output becomes dead and can be released.

Algorithm 4.2: Consumer-count based intermediate release

Input: executed node N , graph consumer count table C

Output: graph with dead temporary arrays cleared

```

1: for each source node  $S$  that feeds  $N$  do
2:    $C[S] = C[S] - 1$ 
3:   if  $C[S] == 0$  and  $S$  is releasable then
4:     clear temporary data stored by  $S$ 
5:   end if

```

```
6: end for
7: if N is a model node and fit has completed then
8:     clear the training input stored only for fitting N
9: end if
```

Several smaller optimizations support this lifetime policy. Cache execution is optional and disabled by default using `cache=False`, because retaining node data can increase memory growth on large datasets. Model nodes release their training arrays after fitting because the fitted estimator is sufficient for inference. The backend also separates output-container allocation from input copying. For preprocessing, the helper `shouldCopyPreprocessInput` decides whether a defensive copy is needed. Scikit-learn transformers normally return transformed data without mutating the original input, so unnecessary copies can be skipped for those objects. Custom functions remain protected when mutation is possible.

4.3.5 Saving and Loading Pipelines

Pipeline objects can be saved as pickle files. Native graph pickling is provided through the PyBind binding for `graph.Graph`, where C++ exposes a `py::pickle` state pair. During `pickle.dump`, Python asks the bound object for its registered pickle state, and the C++ graph returns a Python dictionary that describes the graph structure and stored objects. During load, the inverse state function rebuilds the graph object.

Custom preprocessing functions are handled by storing source code when normal pickle cannot serialize the function object directly. The save path extracts a top-level function or simple lambda source, and the load path executes that source inside a controlled namespace to recover the callable. Functions with closures are rejected because their external captured variables would not be available from source text alone.

4.4 Code Parallelizer Module

4.4.1 Module Responsibility

The Code Parallelizer transforms supported loop-heavy source code into OpenMP-annotated C++. The main Python orchestration file is `accelera/src/utils/parallelizer.py`. Python input is first lowered by `accelera/src/utils/py2cpp_converter.py`. Native loop extraction and pragma insertion are implemented in C++ under `src/ast/` and `src/utils/code_parallelizer_utils.cpp`. The module can process files or in-memory code strings.

The workflow has four stages:

1. Convert supported Python code to C++ when the input is Python.
2. Use Clang-based analysis to extract loops and loop metadata.
3. Predict the loop class using a classifier and rule checks.

4. Insert OpenMP pragmas and return the transformed C++ code.

Figure 4.3 shows the complete workflow of the Code Parallelizer module. The process starts by receiving either Python or C/C++ source code. If the input is Python, it is first converted into a restricted C++ representation. The generated or original C++ code is then parsed and analyzed to extract loop-level features such as nesting, assignments, array accesses, dependencies, and reduction patterns. These features are passed to the trained classifier, while an additional rule-based validation layer checks whether the predicted pragma is safe to apply. Finally, the module inserts the appropriate OpenMP directive and returns the optimized C++ code.




Figure 4.3: Code Parallelizer module workflow

4.4.2 Loop Analysis and Pragma Selection

Loop extraction records syntactic and semantic features such as loop nesting, array access, assignments, reduction patterns, loop-carried dependencies, function calls, and scalar updates. These features are passed to the classifier service. The classifier predicts one of three classes: `none`, `parallel_for`, or `reduction`. The rule layer then validates the predicted class before emitting the pragma.

Algorithm 4.3: Loop classification and OpenMP annotation

Input: source program P

Output: OpenMP-annotated C++ program P'

```

1: if  $P$  is Python then
2:     convert the supported Python subset to C++
3: end if
4: Parse C++ source using the Clang frontend action
5: for each extracted loop  $L$  do
6:     compute structural and dependency features  $F(L)$ 
7:     request class  $c$  from the OpenMP classifier
8:     if  $c$  is parallel_for and dependency checks pass then
9:         insert "#pragma omp parallel for"
10:    else if  $c$  is reduction and reduction variable is valid then
11:        insert "#pragma omp parallel for reduction(...)"
12:    else
13:        keep the loop sequential
14:    end if
15: end for
    
```

16: Return the rewritten C++ source

4.4.3 Classifier Model and Rule-Based Decision

The classifier assets are stored under `models/open_mp_classifier/`. The service loads `model.pkl` and `label_encoder.pkl`, exposes a FastAPI prediction endpoint, and maps loop features to one of the supported pragma classes. The trained model is an `XGBClassifier` over a fixed 40-feature loop representation.

An earlier generator trial attempted to generate pragma decisions using a neural sequence model. The model used a custom PyTorch encoder-decoder architecture: a bidirectional LSTM encoder, an attention-based LSTM decoder, teacher forcing during training, cross-entropy loss, Adam optimization, gradient clipping, and a custom BPE tokenizer. This approach was rejected for two scientific reasons. First, the available dataset contained about 15,000 samples, which was not enough for a sequence model to generalize reliably to unseen program structures; the model overfit easily. Second, generating extra training programs with AI tools risked injecting generator-specific bias, which would lower generalization on real user code. The final approach therefore uses explicit loop features, a classifier, and rule-based safety checks.

4.4.4 Python-to-C++ Converter Support

The Python converter intentionally supports a restricted subset. This makes the generated C++ predictable enough for the parallelizer and avoids silently miscompiling dynamic Python behavior.

Table 4.2: Supported operations in the Python-to-C++ converter

Category	Supported forms	Notes
Values and names	Numeric constants, names, simple lists	Dynamic objects are outside the subset.
Arithmetic	<code>+</code> , <code>-</code> , <code>*</code> , <code>/</code> , <code>%</code> , <code>**</code>	Power maps to <code>std::pow</code> .
Boolean logic	<code>and</code> , <code>or</code> , <code>not</code>	Emitted as C++ boolean expressions.
Comparisons	<code><</code> , <code><=</code> , <code>></code> , <code>>=</code> , <code>==</code> , <code>!=</code>	Chained comparisons are restricted.
Calls	Supported built-ins and direct calls	Unsupported calls raise converter errors.
Indexing	Array-style subscripting	Slice semantics are not supported.
Assignment	Simple assignment and selected augmented assignment	Used for loop updates and reductions.
Control flow	<code>if/else</code> , <code>return</code> , <code>for</code> <code>range(...)</code>	Main loop form used by the parallelizer.

Category	Supported forms	Notes
Functions	Plain function definitions	Decorators, closures, and dynamic dispatch are outside the safe subset.

4.4.5 Integration with Accelera Pipe

The parallelizer is integrated into Accelera Pipe for custom preprocessing functions and custom transformer instances. If a preprocessing function is a user-defined Python function, the pipeline attempts to optimize it through the parallelizer. If the object is a custom transformer, the transform method can be optimized and attached back to the instance. This allows the user to write a normal Python preprocessing step while the system compiles the loop-heavy part into a Python-callable C++ extension.




Figure 4.4: Integration flow for compiled preprocessing inside Accelera Pipe

The integration is applied during graph construction rather than as a separate user-facing execution stage. A custom preprocessing method is inspected, lowered to the supported C++ subset, analyzed for loop-level parallelism, compiled into a Python-callable native module, and then stored back inside the preprocessing node. As a result, the graph still executes a normal preprocessing node, but the callable held by that node may be a compiled OpenMP implementation instead of the original Python loop.

4.5 Auto Preprocessing Module

The Auto Preprocessing module is responsible for automating the data understanding, exploratory data analysis, and preprocessing stages of the data science workflow. The module was designed to reduce repeated manual work and to prepare different types of datasets before model training.




Figure 4.5: Tabular Preprocessing Workflow




Figure 4.6: Text Preprocessing Workflow




Figure 4.7: Image Preprocessing Workflow

4.5.1 Functional Description

The main function of this module is to receive a dataset and user configuration, then return processed data that is ready for machine learning or deep learning models. The module supports tabular, text, and image datasets. The module also generates a report when this option is enabled. The report summarizes information about the dataset and the preprocessing steps applied.

4.5.1.1 Tabular datasets

For tabular datasets it supports both classification and regression tasks. It analyzes the dataset, removes duplicate records, splits the data into training and validation sets, detects column types, handles missing values, encodes categorical features, scales numerical features, applies target preprocessing, performs feature selection, and saves the preprocessing objects needed later for test data.

4.5.1.2 Text datasets

It supports classification tasks. It cleans the text, tokenizes it, applies text preprocessing steps, and converts the text into a numerical representation suitable for machine learning models.

4.5.1.3 Image datasets

It supports image classification by reading images from class folders, validating image files, resizing images, applying augmentation, normalizing pixel values, and preparing the data loaders needed for training.

4.5.2 Modular Decomposition

The Auto Preprocessing module is divided based on the type of input data. The system supports three main types of data: tabular data, text data, and image data. Each type has its own preprocessing workflow and functions because every data type has a different structure and needs different techniques.

4.5.2.1 Tabular Data Preprocessing Functions

`__init__()`: This is the constructor that receives the dataset and the user settings. It validates these settings before starting the preprocessing process.

`handle_bool_types()`: This function handles columns that have bool data type. It finds the columns that have boolean values and converts them into integer values.

`data_overview()`: This function gives a general overview of the dataset before preprocessing.

`drop_target_nulls()`: This function drops rows have null value in the target.

`drop_duplicates()`: This function removes repeated rows from the dataset.

`split_data()`: This function splits the dataset into training and validation sets. The split is based on the validation size and random state selected by the user. The output of this function is `X_train`, `X_val`, `y_train`, and `y_val`.

`get_data_info()`: This function studies each column in the training data. It collects important information such as data type, number of unique values, missing values, and missing percentage. For numerical columns, it also calculates the mean, and median.

`is_drop_column()`: This function decides whether a column should be removed or kept. A column is removed if it is included in the user selected drop list, if it has only one constant value, or if its missing value percentage is higher than the selected threshold.

`drop_col()`: This function removes the selected columns from both the training and validation datasets. It also saves the dropped columns and their reasons in a file. This helps the system apply the same dropping step later on testing or new data.

`detect_column_types()`: This function detects the type of each feature. This step is important because each feature type needs a different preprocessing method. The features are divided into binary columns, ordinal columns, numerical columns, continuous columns, low cardinality categorical columns, high cardinality categorical columns, and unsupported columns. The function also stores the preprocessing steps for each column to be shown later in the report.

Table 4.3: Detected Column Types in Tabular Data

Feature Type	Description	Examples
Binary	Columns containing exactly two unique values.	yes/no, 1/0
Ordinal	Integer columns with a small number of unique values less than or equal to <code>max_ordinal_unique</code> .	Ratings or levels such as 1, 2, 3, 4
Numerical	Integer columns with a number of unique values greater than <code>max_ordinal_unique</code> .	Age, number of transactions
Continuous	Floating-point columns.	Height = 175.5, price = 19.99, temperature = 36.6
Low Cardinality Categorical	Object or string columns with a number of unique values less than or equal to <code>cardinality_threshold</code> .	Product categories, gender, country group
High Cardinality Categorical	Object or string columns with a number of unique values greater than <code>cardinality_threshold</code> .	Product names, cities
Other	Columns with unsupported or non-standard data types. These columns are excluded from the preprocessing pipeline.	Dates, complex objects, unsupported formats

`make_graphs()`: This function creates graphs when the report option is enabled. The type of graph depends on the feature type and the problem type.

`features_preprocessing()`: This function applies the preprocessing steps to the input features. It creates different pipelines based on the detected feature types:

- Numerical and continuous features are filled using the median and scaled using robust scaling.
- Binary features are filled using the most frequent value and encoded using ordinal encoding.
- Low cardinality categorical features are filled using the most frequent value and encoded using binary encoding.
- High cardinality categorical features are filled using the most frequent value and encoded using frequency encoding.
- Ordinal features are filled using the most frequent value and scaled using robust scaling.

These pipelines are combined using `ColumnTransformer`. The transformer is fitted on the training data and then applied to the validation data. The fitted preprocessor and the generated feature names are saved for later use.

`target_preprocessing()`: This function prepares the target column based on the problem type. For classification label encoding is used to convert class labels into numbers. For regression robust scaling is applied. The target preprocessor is saved so the same transformation can be used later.

`features_importance()`: This function selects the most useful features using mutual information. Only features with scores greater than or equal to the selected threshold are kept. The selected features are saved in a file to use it in the inference step.

`common_preprocessing()`: This is the main function that controls the full tabular preprocessing workflow. It calls the previous functions in order. It starts by handling boolean columns, generating the data overview, drop rows have null target, removing duplicates, and splitting the data. Then, it extracts column information, drops unsuitable columns, detects column types, generates graphs, preprocesses the features, preprocesses the target, applies feature selection, and generates the report if enabled. It returns the processed training and validation data.

4.5.2.2 Text Data Preprocessing Functions

`__init__()`: This is the constructor that receives the dataset and the user settings. It validates these settings before starting the preprocessing process. Also removes any columns that are not the text column or the target column. This keeps only the columns needed for text classification. It also downloads the required NLTK resources, such as punkt and stopwords.

`make_graphs()`: This function creates graphs if the report option is enabled. It creates a graph for the text column. It also creates a graph for the target column. These graphs help the user understand the text data and the target class distribution.

`custom_text_tokenizer()`: This function is used to clean the text and split it into words. It converts the text to lowercase, removes extra spaces and unnecessary symbols, and handles some contractions such as `can't` and `won't`. It also separates numbers from letters. Then, the text is tokenized into words. Stopwords are removed, but words like `not` and `no` are kept because they affect the meaning. After that, stemming is used to reduce words to their basic form.

`features_preprocessing()`: This function applies the preprocessing steps to the text column. It fills missing text values in the training and validation data with empty strings. Then it uses the custom tokenizer. It also stores examples of the text

before and after cleaning in the report. Then `TfidfVectorizer` is used to convert the cleaned text into numerical features. The vectorizer uses the user settings, such as the maximum number of features, n-gram range, maximum document frequency, and minimum document frequency. The vectorizer is fitted on the training data and applied to the validation data. The fitted vectorizer is saved as `training_preprocessor.pkl` so the same text transformation can be reused later on testing or new data.

`target_preprocessing()`: This function prepares the target column for classification this function applies label encoding to the target values. The fitted label encoder is saved as `target_preprocessor.pkl`. Also, the text column and target column information are saved in `data_info.pkl`. This helps the system apply the same target transformation later on testing or new data.

`common_preprocessing()`: This is the main function that controls the full text preprocessing workflow. It calls the other functions in the correct order. It creates the data overview. Then, it removes rows that have missing target values and removes duplicate rows. Then, it splits the dataset into training and validation sets. If the report option is enabled, graphs are created. Then, feature preprocessing is applied to clean and vectorize the text. Target preprocessing is applied to encode the class labels. Generates the report if enabled. It returns the processed training and validation data.

4.5.2.3 Image Data Preprocessing

It is used for image classification datasets. In this module, the images are organized inside folders, where each folder represents one class. It is divided into a set of functions, and each function has a specific role in the image preprocessing workflow. The main functions are described as follows:

`__init__()`: This is the constructor that receives the dataset and the user settings. It validates these settings before starting the preprocessing process. It also prepares empty lists to store invalid images and their labels. It saves the image size information in a file so it can be reused later.

`get_classes_mapping()`: This function creates the class mappings. It gives each class name a numerical label. It creates two mappings: one from class name to label, and another from label to class name. These mappings are saved in files. This helps the system use the same labels during training, validation, testing, and prediction.

`data_preparing()`: This function reads the images from the class folders. It goes through each class folder and checks each image path. If the image is valid and can be opened, its path is added to the list of valid image paths, and its class

label is added to the labels list. If the image is invalid or corrupted, it is added to the invalid images list with its label. If there are no valid images, the function raises an error because the preprocessing process cannot continue without valid image data.

`data_overview()`: This function collects general information about the image dataset. For the training folder, it stores the class names, number of valid images, random sample images, number of invalid images, invalid image paths, and the class mapping. If a validation folder is provided, it also stores the same type of information for the validation folder. This information is used later in the final preprocessing report.

`make_graphs_label_summary()`: This function creates graphs that show the label distribution. It creates a label summary graph for the training folder. If a validation folder is provided, it also creates a label summary graph for the validation folder.

If the training data is split into training and validation sets, the function also creates graphs for the labels after splitting. These graphs help the user understand if the classes are balanced or not.

`make_graphs_sample()`: This function creates sample image graphs. It shows random samples from the training folder. If a validation folder is provided, it also shows random samples from the validation folder.

If the training data is split into training and validation sets, it also shows sample images after splitting. This helps the user visually check the image dataset.

`make_graphs_loader()`: This function creates sample images after using the `DataLoader`. It takes a small number of images from the training `DataLoader` and displays them.

If a validation `DataLoader` exists, it also displays sample images from it. This step helps confirm that the images are loaded correctly after resizing, normalization, and augmentation.

`make_graphs()`: This function controls all graph generation steps. It prepares the graph section in the report, then calls the label summary graphs, sample image graphs, and `DataLoader` sample graphs.

`get_loaders()`: This function creates the training and validation datasets and `DataLoaders`. First, it creates a `ClassificationImageDataset` for the training images using the image paths, labels, image size, and augmentation settings.

Then, it creates the training `DataLoader` with the selected batch size and uses shuffling. Shuffling is used for training because it helps the model see the data in a different order.

If validation data exists, the function creates a validation dataset and validation DataLoader. Augmentation is not applied to the validation data, and shuffling is not used because validation should be more stable.

`common_preprocessing()`: This is the main function that controls the full image pre-processing workflow. It starts by creating the class mappings. Then, it prepares the training images by collecting valid image paths and labels and saving invalid images.

If a validation folder is provided, it also prepares the validation images. After that, it creates the data overview, splits the data if needed, creates the DataLoaders, generates graphs, and creates the final image preprocessing report.

At the end, the function returns the training DataLoader and validation DataLoader. These outputs are ready to be used for training an image classification model.

4.5.2.4 Testing and Reusability:

For all supported data types, the saved preprocessing objects are reused during testing. This is important because the testing data must be processed in the same way as the training data.

- For tabular data, the system loads the saved preprocessors, selected features, dropped columns, and target preprocessor.
- For text data, the system loads the saved TF-IDF vectorizer and label encoder.
- For image data, the system loads the saved class mappings and image size.

This makes the preprocessing module reusable and consistent for training, validation, testing, and future prediction.

4.5.3 Design Constraints

Several constraints affected the design of the Auto Preprocessing module:

- The module must avoid data leakage. For this reason, preprocessing objects are fitted only on the training data and then applied to validation or testing data.
- The module must support different data types. Therefore, separate preprocessing paths were designed for tabular, text, and image datasets.
- The tabular path assumes that the data is structured in rows and columns and that the task is either classification or regression.
- The text path assumes that the dataset contains a text column and a target label for a text classification task.

- The image path assumes that image classification data follows a folder-based class structure, where each folder represents one class.
- The module must keep preprocessing reproducible. Therefore, fitted preprocessing objects and selected features are saved for later use.
- The module must be general enough for different datasets, so the design uses rule-based decisions instead of manual preprocessing code for each dataset.

4.5.4 Other Description of the Module

The Auto Preprocessing module is not only designed to transform data. It is also designed to help the user understand the dataset before and after preprocessing. This is achieved through the generated report and visualizations. For tabular data, the report can show missing-value statistics, duplicate information, feature distributions, target distribution, correlation matrix, removed columns, preprocessing decisions, and previews of the final processed training and validation data.

Another important design point is that the module separates training preprocessing from testing preprocessing. During training, the module fits the preprocessing objects and saves them. During testing, the saved objects are loaded and applied directly, which ensures that new data is transformed in the same way as the training data.

4.6 AutoML Module

4.7 Deployment Module

4.8 Benchmark Platform Module

The Benchmark Platform module is designed to provide a shared environment where users can create machine learning benchmarks, submit predicted results, and compare submissions using selected evaluation metrics. The module includes a backend API, database schemas, Python validation and scoring scripts, and a frontend interface for users and admins.

4.8.1 Functional Description

The main function of this module is to support fair comparison between different machine learning solutions. A user can create a benchmark by providing the benchmark title, description, problem type, target column, dataset link, test dataset link, ground-truth target link, and evaluation metric. The system validates the provided links and

checks that the test set does not contain the target column while the target file contains the expected target column.

After a benchmark is created, other users can submit their prediction file and repository link. The system validates the submitted prediction file, calculates the score using the selected Scikit-learn metric, stores the submission, and displays submissions on a leaderboard. The leaderboard is ordered based on the metric configuration, where some metrics are better when their value is higher and others are better when their value is lower.

4.8.2 Modular Decomposition

The Benchmark Platform is decomposed into several smaller components:

4.8.2.1 Benchmark Management

This component handles benchmark creation, listing, filtering, viewing, and deletion. It stores benchmark information such as title, description, problem type, dataset links, target column, selected metric, metric parameters, creation date, and creator.

4.8.2.2 Metric Management

This component allows admin users to define evaluation metrics. Each metric contains a display name, the corresponding Scikit-learn metric function name, needed parameters, problem type, and whether a higher or lower score is better. Metrics are linked to benchmarks and used during submission scoring.

4.8.2.3 Submission Management

This component allows users to submit prediction files for a specific benchmark. Each submission stores the benchmark id, submitting user, repository link, prediction file link, score, and submission date. Users can also update their existing prediction file or delete their submission if they have permission.

4.8.2.4 Validation and Scoring Scripts

The backend calls Python scripts to validate benchmark files and score submissions. During benchmark creation, the validation script checks that the test file does not include the target column and that the target file contains the required target column. During submission, the scoring script loads the ground-truth file and the submitted prediction file, checks their length and target column, then computes the selected Scikit-learn metric.

4.8.2.5 Leaderboard Interface

The frontend displays benchmark submissions in a leaderboard. It shows the submitter information, repository link, score, and submission date. The sorting direction depends on the selected metric, so the platform can support both metrics where higher is better and metrics where lower is better.

4.8.3 Design Constraints

Several constraints affected the design of the Benchmark Platform:

- The module currently supports classification and regression problem types only.
- Dataset, test, target, and prediction files are provided using links, mainly Google Drive file links, to reduce server storage requirements.
- The test dataset must not contain the target column. This prevents data leakage and makes the evaluation fair.
- The target file and submitted prediction file must contain the selected target column and must have the same number of rows.
- Evaluation metrics must be defined using valid Scikit-learn metric names because the scoring script calls Scikit-learn metrics dynamically.
- Only authenticated users can create benchmarks or submit predictions. Admin permission is required for metric creation.
- A user is allowed to create only one active submission for the same benchmark. If the user wants to change the result, the existing submission is updated instead of creating duplicate submissions.

4.8.4 Other Description of the Module

The Benchmark Platform separates benchmark creation from submission evaluation. This makes the workflow clearer: first, the benchmark creator defines the task and evaluation rules; then, users submit predictions under the same conditions. This separation helps ensure that all users are evaluated using the same ground-truth file and the same metric configuration.

Another important design point is the use of Python scripts for file validation and metric scoring. The backend is implemented using Node.js and Express, while the scoring process uses Python because it depends on Scikit-learn metrics and dataset loading utilities already used in the project. This design allows the web platform to manage users, benchmarks, and submissions while Python handles the machine learning evaluation part.

Chapter 5: System Testing and Verification

This chapter explains how the project was tested to make sure that the system works correctly. The testing was done in different ways because not all parts of the project can be tested using the same method. Some parts were tested using unit tests, some parts were checked manually, and some results were compared with other tools or expected outputs. The main purpose of testing was to make sure that the system can take the input, apply the required steps, save the needed files, and return the correct output without errors.

5.1 Testing Setup

The testing was done in a Python environment. The main tools used in testing were Python, Pytest, Pandas, NumPy, Scikit-learn, and the other libraries used in the project implementation. The system was tested using sample datasets and different test cases. The output of the system was checked to make sure that it is correct and in the expected format to make sure that they were created correctly. Some tests were written and run automatically using Pytest. Other tests were checked manually by comparing the actual output with the expected output. This testing setup helped make sure that the project works correctly.

5.2 Testing Plan and Strategy

The testing plan was based on testing the project step by step. First, the main parts of the system were tested separately to make sure that each part works correctly. After that, the complete system was tested as one workflow to make sure that all parts work together without errors. Different testing methods were used depending on the part being tested. Some parts were tested using unit tests with Pytest. Other parts were checked manually by comparing the actual output with the expected output. In some cases, the results were also compared with other tools to make sure that the system gives correct and reasonable results.

5.2.1 Module Testing

5.3 Accelerera Pipe Module

The Accelerera Pipe experiments compare three implementations under the same train, validation, and test split. The first implementation is a manually written scikit-learn

loop over the candidate pipelines. The second uses `sklearn.pipeline.Pipeline` to represent the same candidates. The third uses Accelera Pipe with graph branches. Each candidate is selected using validation accuracy, and the final selected path is evaluated on the test set. This protocol checks whether the graph execution improves latency without changing the selected model or the measured predictive accuracy.

Table 5.1: Accelera Pipe latency scaling in seconds

Samples	sklearn	sklearn Pipeline	Accelera Pipe
1,000	0.06	0.05	0.34
5,000	2.47	2.12	2.02
10,000	2.48	2.48	1.79
50,000	23.96	23.25	14.81
100,000	77.79	77.10	67.04
250,000	803.56	803.01	494.64

Table 5.1 shows the main performance trend without compressing the font. At very small sizes, Accelera Pipe is slower because the native graph setup, branch creation, and Python/C++ boundary cost dominate the actual model computation. Once the dataset grows, those fixed costs become small compared with fitting multiple candidate branches. At 50,000 samples, Accelera Pipe reduces runtime from 23.96 seconds to 14.81 seconds. At 250,000 samples, it reduces runtime from about 803 seconds to about 495 seconds, saving more than five minutes while selecting the same candidate.

Table 5.2: Accelera Pipe RSS memory scaling in MB

Samples	sklearn	sklearn Pipeline	Accelera Pipe
1,000	1.62	0.16	19.32
5,000	3.62	1.09	21.62
10,000	11.47	2.19	7.25
50,000	126.95	9.75	209.25
100,000	139.70	47.65	300.93
250,000	303.95	28.52	594.86

Table 5.2 shows the current cost of graph-level parallelism. Accelera Pipe uses more RSS memory because multiple branch states and intermediate arrays can exist at the same time. The memory optimizations in the methodology section, especially consumer-count based release, duplicate first-node sharing, and guarded preprocessing copies, target exactly this issue.

Table 5.3: Selected candidate and accuracy across sample sizes

Samples	Selected candidate	Validation accuracy	Test accuracy
1,000	MinMaxScaler → SVC	0.7900	0.8250
5,000	MinMaxScaler → SVC	0.9160	0.9100
10,000	MinMaxScaler → SVC	0.9385	0.9260
50,000	MinMaxScaler → SVC	0.9565	0.9598
100,000	StandardScaler → SVC	0.9675	0.9709
250,000	StandardScaler → SVC	0.9774	0.9768

Table 5.3 confirms that the speedup does not come from changing the search space. Accelera Pipe evaluates the same candidate pipelines and reports the same validation and test behavior. The selected preprocessing strategy changes from Min-Max scaling to standard scaling only when the larger dataset makes StandardScaler with SVC the best validation candidate. This is expected because the candidate set and selection rule are fixed while the data distribution becomes better represented at larger sizes.

Table 5.4: Largest Accelera Pipe comparison

Implementation	Time (s)	RSS MB	VMS MB	Test accuracy
sklearn manual loop	803.56	303.95	315.41	0.9768
sklearn Pipeline	803.01	28.52	28.61	0.9768
Accelera Pipe	494.64	594.86	787.47	0.9768

Table 5.4 isolates the largest run: 150,000 training samples, 50,000 validation samples, and 50,000 test samples. Accelera Pipe keeps exactly the same selected path, StandardScaler followed by SVC, and the same test accuracy. The runtime drop is 308.91 seconds relative to the manual scikit-learn loop and 308.36 seconds relative to scikit-learn Pipeline. This result is important because it demonstrates the value of graph scheduling on realistic branch-heavy workloads. The memory columns also show the next engineering target: scikit-learn Pipeline is extremely memory efficient, while Accelera Pipe trades memory for parallel branch execution and native graph state.

5.4 Code Parallelizer Module

The OpenMP classifier was evaluated as a three-class loop classifier. The labels are `none`, `parallel_for`, and `reduction`. The classification report shows that the model is balanced enough to guide the rule layer, while the final pragma decision remains protected by dependency checks.

Table 5.5: OpenMP classifier report

Class	Precision	Recall	F1-score	Support
none	0.82	0.85	0.84	3048
parallel_for	0.86	0.81	0.83	2696
reduction	0.79	0.82	0.80	775
accuracy			0.83	6519
macro avg	0.82	0.83	0.83	6519
weighted avg	0.83	0.83	0.83	6519

Table 5.5 should be read as a guidance-quality result, not as a proof that every predicted loop is safe. The classifier is responsible for identifying likely parallelization opportunities. The rule layer is responsible for rejecting unsafe or unsupported cases. This two-stage design is useful because source code parallelization has correctness constraints that pure statistical classification should not decide alone.

Table 5.6: Hard parallelizer evaluation

File	Baseline ms	Parallel ms	Speedup	Output match
hard_eval_000.cpp	1786.28	195.47	9.138×	true
hard_eval_001.cpp	63.98	31.47	2.033×	true
hard_eval_002.py	42.66	7.67	5.563×	true

Table 5.6 evaluates harder programs from `data/hard_test_files`. The benchmark compiles and runs both the baseline and the generated parallelized code, then compares their outputs. The first file obtains the largest improvement because its loop body has enough work to amortize OpenMP thread overhead. The second file still improves, but the smaller absolute runtime limits the maximum speedup. The Python-origin file demonstrates the full converter path: Python subset to C++, then loop extraction, then pragma insertion. Across the three files, the evaluation extracted 11 loops, generated 8 pragmas, matched all 8 expected pragmas, reached average similarity 1.0, and obtained an average speedup of 5.578×

Table 5.7: Linux parallelizer and Accelera Pipe integration benchmark

Mode	Samples	Time (s)
Python custom preprocessing function	500,000	6.83
End-to-end optimized run, first measured process	500,000	5.88
End-to-end optimized run, warmed method and module cache	500,000	0.08

Table 5.7 reports the Linux direct `examples/parallel_accpipe.py` run after normalizing the example execution directory to the repository root. The Python imple-

mentation takes 6.83 seconds on 500,000 samples. The first optimized process in the measurement sequence takes 5.88 seconds because it still pays process-local setup and cache lookup costs. The immediate repeated run takes 0.08 seconds after the Python-method cache and compiled-module cache are both warm. The validation in the example confirmed that the optimized output matched the Python output. This result shows why the integration needs two cache layers: one layer avoids repeating Python-to-C++ conversion and OpenMP insertion, while the lower compiled-module cache avoids recompiling the pybind extension.

Table 5.8: Linux direct example-script verification runs

Example	Baseline time (s)	Optimized time (s)	Validation
accpipe_demo	2.48	1.79	same accuracy, 0.9260
parallel_accpipe, run 1	6.83	5.88	outputs match
parallel_accpipe, run 2	6.83	0.08	outputs match
code_parallelizer, run 1	5.81	0.014	outputs match
code_parallelizer, run 2	6.05	0.015	outputs match
save_load, run 1	1.89	2.69 / 2.65	same accuracy, 0.268
save_load, run 2	2.53	2.20 / 2.14	same accuracy, 0.268

Table 5.8 summarizes the Linux examples invoked by `examples/Makefile`, but each file was run directly rather than through `Make`. Cache-sensitive files were repeated. The standalone `code_parallelizer_demo.py` result measures Python script execution against the generated OpenMP executable for `hard_eval_002.py`; the C path also reported about 0.26–0.27 seconds of compilation time, outside the listed executable runtime. The save/load example verifies persistence rather than pure acceleration: both `sklearn` and `Accelera` load persisted state and produce the same accuracy, while the second `Accelera` run is slightly faster than the `sklearn` baseline.

Table 5.9: Windows direct example correctness and path timing

Example	Run/cache state	Baseline (s)	Accelera/C path (s)	Validation
accpipe_demo	normal	1.75 / 1.78	1.03	same test accuracy, 0.9260
parallel_accpipe	isolated cache, run 1	6.66	6.00	outputs match
parallel_accpipe	isolated cache, run 2	6.58	0.09	outputs match
code_parallelizer	isolated cache, run 1	5.59	1.04	outputs match
code_parallelizer	isolated cache, run 2	5.61	1.05	outputs match
save_load	normal	0.57	0.51 / 0.53	same accuracy, 0.268

Table 5.10: Windows compilation and process overhead

Example	Run/cache state	Compile/link (s)	Wall time (s)	Interpretation
accpipe_demo	normal	–	6.53	graph example startup
parallel_accpipe	isolated cache, run 1	included	14.37	cold compile path
parallel_accpipe	isolated cache, run 2	cached	8.33	warm cache path
code_parallelizer	isolated cache, run 1	0.87	8.74	temporary executable
code_parallelizer	isolated cache, run 2	1.04	9.00	temporary executable
save_load	normal	–	2.71	persistence example

Tables 5.9 and 5.10 report the same Makefile example set on Windows. The first table keeps the user-visible comparison: baseline runtime, optimized path runtime, and validation result. The second table separates compilation and process-level overhead so the timing table remains readable. The examples were run one by one in Makefile order. The cache-sensitive examples were then repeated under an isolated Accelera cache so that the first run represents cold cache behavior and the second run represents warm cache behavior. The strongest cache effect appears in `parallel_accpipe.py`: the cold end-to-end Accelera path takes 6.00 seconds because it must convert the function, select loops, compile a pybind extension, link it, and load it into Python. The immediate warm-cache run takes 0.09 seconds because the processed-code cache, method cache, and compiled-module cache avoid repeating those setup stages.

The Windows `code_parallelizer_demo.py` runs show a different limitation. The generated C/OpenMP executable is compiled into a temporary directory on every run, so the second run does not reuse the final linked program. The measured C executable subprocess takes about 1.04 seconds, while the instrumented main-body timer inside the generated program reports only about 0.024–0.025 seconds. This gap is Windows process startup, OpenMP runtime initialization, dynamic library loading, and executable loading overhead. The compile/link step itself adds about 0.87–1.04 seconds. On Linux, the same benchmark has lower launch and linking overhead in the reported measurements, so executable runtime is closer to the timed main-body work. For this reason, Windows results should be interpreted as end-to-end toolchain latency results, not only as loop-kernel speed results.

5.4.0.1 Auto Preprocessing Module

Most of the functions were tested using unit tests. These tests were used to make sure that the functions give the expected output and do not produce errors. Real datasets were also used to test the system in a more practical way. These datasets were collected from Kaggle for different data types. For all datasets, the data was split using a test size of 0.2 and a random state of 42. The model performance was then measured on the test data. For tabular data, both classification and regression datasets were used. After applying preprocessing, the output was passed to a Random Forest model

to check that the processed data was ready for training and testing. For text data, the processed output was passed to an `XGBClassifier`. This helped check that the text preprocessing steps converted the text into numerical features correctly. For image data, the output of the image preprocessing module was passed to a ResNet model. This helped check that the image data loaders were created correctly and could be used directly for model training. The results showed that the tested modules produced valid outputs and were suitable to be used in the training step.

Table 5.11: Datasets Used in Module Testing

Dataset Name	Shape	Dataset Type	Problem Type
startup_founder_burnout_2026	(50000, 29)	Tabular	Classification
healthy_diet_calorie_intake	(6000, 15)	Tabular	Classification
ai_student	(50000, 16)	Tabular	Classification
gaming_addiction	(250, 49)	Tabular	Classification
student_exam_performance_dataset	(10000, 17)	Tabular	Classification
heart	(1025, 14)	Tabular	Classification
Dataset_spine	(310, 7)	Tabular	Classification
diabetes_dataset	(100000, 31)	Tabular	Classification
student_placement_synthetic	(100000, 18)	Tabular	Classification
Titanic-Dataset	(891, 12)	Tabular	Classification
customer_purchase_data	(1500, 9)	Tabular	Classification
global_ai_jobs	(90000, 35)	Tabular	Regression
CarPrice_Assignment	(205, 26)	Tabular	Regression
Housing	(545, 13)	Tabular	Regression
PetImages	(24998,)	Image	Classification
DailyDialog	(11327, 2)	Text	Classification
TestReviews	(4321, 2)	Text	Classification

Table 5.12: Models and Evaluation Metrics Used in Module Testing

Dataset Name	Model Used	Evaluation Metric	Accelera Result
startup_founder_burnout_2026	RandomForest	Macro F1-score	0.9999
healthy_diet_calorie_intake	RandomForest	Macro F1-score	1.0000
ai_student	RandomForest	Macro F1-score	0.4799
gaming_addiction	RandomForest	Macro F1-score	1.0000
student_exam_performance_dataset	RandomForest	Macro F1-score	0.8573
heart	RandomForest	Macro F1-score	0.7860
Dataset_spine	RandomForest	Macro F1-score	0.7840
diabetes_dataset	RandomForest	Macro F1-score	0.9985
student_placement_synthetic	RandomForest	Macro F1-score	0.9986
Titanic-Dataset	RandomForest	Macro F1-score	0.7906
customer_purchase_data	RandomForest	Macro F1-score	0.9385
global_ai_jobs	RandomForest	R-squared Score	0.9227
CarPrice_Assignment	RandomForest	R-squared Score	0.9562
Housing	RandomForest	R-squared Score	0.6037
PetImages	ResNet model	Accuracy	0.97
DailyDialog	XGBClassifier	Accuracy	Add result
TestReviews	XGBClassifier	Accuracy	Add result

5.4.0.2 AutoML Module

5.4.0.3 Deployment Module

5.4.0.4 Benchmark Platform Module

5.4.1 Integration Testing

5.4.2 Testing Schedule

5.4.3 Comparative Results to Previous Work

Chapter 6: Conclusions and Future Work

6.1 Faced Challenges

The main implementation challenges were preserving correctness while sharing branch computations, releasing intermediate arrays only after their final consumer, and limiting generated OpenMP code to loops that pass dependency and conversion checks. The experiments also identify memory usage during branch-level parallel execution as the main remaining graph-backend challenge.

6.2 Gained Experience

6.3 Conclusions

Accelerera Pipe provides a graph-based execution model for branch-heavy machine-learning pipelines and supports fitted executed graphs. The Code Parallelizer provides a source-to-source path from supported Python or C++ loops to OpenMP-annotated C++ using converter rules, loop feature extraction, classifier guidance, and safety checks. The available experiments show runtime improvements on large branch-selection workloads while preserving selected models, accuracy, and output equivalence.

6.4 Future Work

Future work should continue reducing graph memory consumption while preserving parallel branch execution, and should extend safe converter and loop-analysis coverage without weakening correctness checks.

References

Bibliography

- [1] "Data Preprocessing and Feature Engineering for Data Mining: Techniques, Tools, and Best Practices." Available at: <https://www.proquest.com/openview/17dfa731b6147c8939e3c5cc784897ce/1?pq-origsite=gscholar&cbl=5046920>

Chapter A: Development Platforms and Tools

The implemented modules use Python, C++, PyBind11, CMake, Clang-based AST analysis, OpenMP, NumPy, and Scikit-learn. The available experiments include Linux and Windows execution paths.

Chapter B: Use Cases

Chapter C: User Guide

Chapter D: Code Documentation

Chapter E: Feasibility Study