



Cairo University
Faculty of Engineering
Department of Computer Engineering

Accelera



A Graduation Project Report Submitted
to
Faculty of Engineering, Cairo University
in Partial Fulfillment of the Requirements for the Degree
of
Bachelor of Science in Computer Engineering.

Presented by

Nesma Osama
9220912

Supervised by

Dr. Salma Abdel Monem

April 17, 2026

All rights reserved. This report may not be reproduced in whole or in part, by photocopying or other means, without the permission of the authors or department.

Abstract

This report explains my individual work in three parts of the Accelera project. The first part is Auto Preprocessing, which prepares tabular, text, and image datasets before model training. The second part is the Accelera Pipeline reporting part, which generates reports for pipeline graphs, branches, and metrics. The third part is the Benchmarking module, which allows prediction submissions to be compared using the same dataset and evaluation metric. In this report, I describe the main idea, the steps I followed, the generated reports, the benchmark workflow, and the testing results.

Table of Contents

List of Figures

List of Tables

List of Algorithms

- 3.1 General tabular preprocessing flow 24
- 3.2 Text tokenizer flow from the implementation 30
- 3.3 Image training preprocessing flow from the implementation 33
- 3.4 Pipeline graph report generation flow 36
- 3.5 Benchmark creation flow 38
- 3.6 Submission scoring flow 39
- A.1 Calling tabular training preprocessing 50
- A.2 Calling tabular regression preprocessing 50
- A.3 Calling tabular testing preprocessing 50
- A.4 Calling text training preprocessing 51
- A.5 Calling text testing preprocessing 51
- A.6 Calling image training preprocessing 52
- A.7 Calling image testing preprocessing 52
- A.8 Calling the pipeline graph report 53

List of Abbreviations

AI	Artificial Intelligence
CSV	Comma-Separated Values
EDA	Exploratory Data Analysis
HTML	HyperText Markup Language
JSON	JavaScript Object Notation
KNN	K-Nearest Neighbors
ML	Machine Learning
NLP	Natural Language Processing
PDF	Portable Document Format
TF-IDF	Term Frequency-Inverse Document Frequency
UI	User Interface

Chapter 1: Introduction

1.1 Motivation and Justification

In many machine learning projects, building the model is not the only important step. Before training the model, the dataset usually needs many preparation steps such as cleaning, checking missing values, splitting the data, encoding categorical columns, scaling numerical values, and making sure there are no errors. These steps are repeated in most projects, so doing them manually every time can take a lot of time and may lead to mistakes.

Accelerera tries to make these steps easier by providing them in one framework. The Auto Preprocessing module is responsible for preparing the data automatically. It applies the needed preprocessing steps and saves the fitted objects, such as encoders and scalers, so the same steps can be used later on new data. This is important because using different preprocessing steps between training and testing can give wrong results.

The Reporting in Pipeline module helps the user understand the pipeline results in a clearer way. It shows the generated graphs, explains the different branches, and allows the user to compare the results of these branches. Instead of only displaying numbers in the terminal, the module creates HTML reports that include graphs, tables, and metric summaries.

The Benchmarking module is used when different users want to compare their results on the same task. The benchmark creator prepares the test data and keeps the correct target values hidden. Then, users submit their predictions, and the system evaluates them using the same metric. This makes the comparison fair because all submissions are tested using the same rules.

These modules help reduce manual work, make the machine learning process more organized, and give users clearer results.

1.2 Project Objectives and Problem Definition

In machine learning projects, preparing the data and understanding the results are important steps before and after model training. However, these steps may become difficult when they are done manually, especially when the dataset contains

different types of data or when the user needs to repeat the same process many times.

This project focuses on improving these steps inside Accelerera. It aims to support the user during data preprocessing, pipeline reporting, and benchmarking. The Auto Preprocessing module prepares different types of datasets for machine learning tasks. The Reporting in Pipeline module presents the pipeline outputs in a clear form using reports, graphs, tables, and metric summaries. The Benchmarking module gives users a fair way to submit their predictions and compare their results using the same evaluation rules.

the main objective of this work is to make the machine learning workflow more organized, easier to understand, and less dependent on manual steps.

1.3 Project Outcomes

The implemented work produces three main outcomes. First, the Auto Preprocessing module prepares tabular, text, and image datasets and saves the fitted artifacts needed for later testing. Second, the reporting code generates readable preprocessing and pipeline outputs using HTML pages, tables, graphs, and metric summaries. Third, the Benchmarking module supports benchmark creation, submission validation, scoring, and leaderboard display.

1.4 Document Organization

Chapter 2 summarizes the related concepts and previous work. Chapter 3 explains the system design and architecture of Auto Preprocessing, Pipeline Reporting, and Benchmarking. Chapter 4 presents testing and comparative evaluation. Chapter 5 concludes the report and lists future improvements. The user guide is in the appendix.

Chapter 2: Literature Survey

In this chapter, I give an overview of the topics, algorithms, and techniques needed to understand my work in the project. I discuss the technical subjects and background knowledge.

2.1 Background and Previous Work

This section introduces the main background needed to understand the Accelera project as a whole. It explains the important concepts, techniques, and tools that were used in different parts of the system. These background topics help clarify how the project supports data preparation, pipeline execution, report generation, model evaluation, and benchmarking. The goal of this section is to give the reader enough context before moving to the methodology and implementation chapters.

2.1.1 Data Science and Machine Learning

Data science is the field that uses data, statistics, programming, and domain knowledge to understand a problem and extract useful insights from it. It is usually used when we have raw data and we want to turn it into useful information or predictions. Machine learning is one of the main parts of data science. Instead of writing fixed rules for every case, we train a model on examples, and the model learns patterns from these examples. After training, the model can use what it learned to make predictions on new data.

In supervised machine learning, the training data contains input features and a target label. The model learns the relationship between the features and the target. If the target is a class, such as spam or not spam, the problem is called classification. If the target is a continuous number, such as price or temperature, the problem is called regression.

2.1.1.1 Data Exploration and Preprocessing [?]

Exploratory Data Analysis, or EDA, is the step where we inspect the dataset before building the model. We summarize columns, draw visualizations, check missing values, and try to understand relationships between variables. This step helps us notice problems early instead of discovering them after training the model.

Data preprocessing is the stage of transforming raw data into a clean and usable format for machine learning models. In real datasets, the data may contain missing values, duplicated rows, different feature scales, and categorical columns that cannot be used directly by many models. Because of this, preprocessing is usually considered an essential part of a machine learning workflow.

Missing value imputation is used when some values in a dataset are empty. Instead of removing all rows that contain missing values, the missing values can be replaced using a suitable value, such as the mean, median, mode, or another estimated value. Robust scaling is a scaling technique that uses the median and the interquartile range instead of the mean and standard deviation. It is useful when the data contains outliers, because the median and interquartile range are less affected by extreme values:

$$x' = \frac{x - \text{median}(X)}{\text{IQR}(X)}$$

where x' is the scaled value, x is the original value, $\text{median}(X)$ is the median of the feature values, and $\text{IQR}(X)$ is the interquartile range.

Categorical encoding is the process of converting text categories into numeric values. Many traditional machine learning models and preprocessing pipelines expect numerical input. Two examples are binary encoding and frequency encoding. Binary encoding first gives each category a number, then converts this number into binary form and stores the binary digits in separate columns. Frequency encoding replaces each category with how often it appears in the dataset:

$$f(c) = \frac{\text{count}(c)}{N}$$

where $f(c)$ is the frequency of category c , $\text{count}(c)$ is the number of rows containing this category, and N is the total number of rows.

2.1.1.2 Model Evaluation Metrics [?]

After training a model, we need metrics to check whether the model is performing well or not. The selected metric depends on the type of problem. Classification models are evaluated using metrics that compare predicted classes with actual classes, while regression models are evaluated using metrics that measure the difference between predicted numerical values and real values.

Classification metrics: Classification metrics are used when the output is a class or category, such as spam or not spam, positive or negative, or one image class from many classes. The confusion matrix is the base idea behind many classification metrics. It compares the actual classes with the predicted classes. For binary classification, it contains four values: true positives, true negatives, false positives, and false negatives.

- **True Positive (TP):** The model predicts positive, and the actual class is positive.
- **True Negative (TN):** The model predicts negative, and the actual class is negative.
- **False Positive (FP):** The model predicts positive, but the actual class is negative.

- **False Negative (FN):** The model predicts negative, but the actual class is positive.

Accuracy measures how many predictions are correct out of all predictions:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP is true positives, TN is true negatives, FP is false positives, and FN is false negatives. Accuracy is simple and useful when the classes are balanced, but it may be misleading when one class is much larger than the other.

Precision measures how many predicted positive samples are actually positive:

$$Precision = \frac{TP}{TP + FP}$$

Precision is useful when false positives should be reduced.

Recall measures how many actual positive samples were correctly detected:

$$Recall = \frac{TP}{TP + FN}$$

Recall is useful when false negatives should be reduced.

The F1-score combines precision and recall:

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

It is useful when both precision and recall are important, especially with imbalanced datasets.

Regression metrics: Regression metrics are used when the output is a numerical value, such as price, score, or temperature.

Mean Absolute Error measures the average absolute difference between the real values and predicted values:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

where y_i is the actual value, $\hat{y}_i$ is the predicted value, and n is the number of samples. MAE is easy to understand because it gives the average error in the same unit as the target value.

Mean Squared Error gives more penalty to large errors:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

This means large mistakes affect the score more than small mistakes.

Root Mean Squared Error is the square root of MSE:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

RMSE is easier to read than MSE because it returns the error to the same unit as the target value.

R-squared score measures how well the regression model explains the variation in the actual values:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

where y_i is the actual value, $\hat{y}_i$ is the predicted value, and $\bar{y}$ is the mean of the actual values. A value close to 1 means the model explains the data well, while a value close to 0 means the model does not explain the data well.

2.1.2 Natural Language Processing (NLP) [?]

Natural Language Processing is the field that helps computers understand and work with human language. It is used in many applications such as text classification, sentiment analysis, spam detection, and chatbots. Text data cannot be used directly by most machine learning models, so it needs to be cleaned and converted into numbers first.

In text preprocessing, the text may be converted to lowercase, cleaned from unnecessary symbols, split into words, and simplified using stemming. After that, methods such as TF-IDF are used to transform the text into numerical features. In Accelera, these steps are used to prepare text datasets so they can be used for machine learning models.

2.1.2.1 Term Frequency-Inverse Document Frequency (TF-IDF)

Term Frequency-Inverse Document Frequency is one common technique used to convert text into numerical features. It gives each word a weight based on how important it is in one document and across the full dataset. Term Frequency measures how often a word appears in a document. Inverse Document Frequency gives lower weight to words that appear in many documents, because these words are usually less useful for classification:

$$TFIDF(t, d) = TF(t, d) \times IDF(t)$$

where $TFIDF(t, d)$ is the weight of term t in document d , $TF(t, d)$ is the frequency of the term in the document, and $IDF(t)$ measures how rare the term is across all documents.

2.1.3 Data Augmentation

Data augmentation is a technique used to increase the diversity of a dataset without actually collecting new data. It works by applying transformations to existing samples to create new modified versions. For example, in image classification, augmentation may include flipping, rotating, cropping, or changing brightness. This can help the model generalize better because it sees more variations during training.

2.2 Comparative Study of Previous Work

The most direct external comparison in this report is AutoCleanML because it also automates common tabular preprocessing tasks. The comparison focuses on the effect of preprocessing on model results, using the same validation split and the same machine learning models where possible.

2.3 Implemented Approach

2.3.1 Auto Preprocessing

In this module I decided to implement an automated preprocessing approach for tabular datasets. The approach was designed to prepare data for machine learning models by applying suitable preprocessing steps automatically according to the type and characteristics of each column and some user configurations.

The general framework of the approach is as follows:

1. Take the dataset and some user configurations.
2. Remove duplicate records.
3. Split the data into training and validation sets.
4. For each column, we get this information: data type, percentage of missing values, unique values, and basic stats.
5. Drop unnecessary columns, such as constant columns or columns with missing values above the selected threshold.
6. Detect column types automatically.
7. Make Exploratory Data Analysis.
8. Apply missing value imputation.
9. Apply categorical encoding.
10. Apply robust scaling.

11. Apply feature selection using mutual information.
12. Automatically saves all needed files for inference.
13. Generate a report.

The chosen approach follows a classical automated preprocessing pipeline. The main preprocessing decisions are:

- **Missing values:** Median imputation is used for numerical features because it is simple and more robust to outliers than mean imputation. For categorical features, most-frequent imputation is used because it replaces missing values with valid existing categories.
- **Numerical scaling:** `RobustScaler` is used for numerical features because tabular datasets may contain outliers or skewed values. It depends on the median and interquartile range, which makes it more stable than standard scaling when extreme values exist.
- **Categorical encoding:** The encoding method is selected according to the feature type and cardinality. Binary columns are encoded using ordinal encoding, low-cardinality categorical columns are encoded using binary encoding, and high-cardinality categorical columns are encoded using frequency encoding. This helps reduce the number of generated features and avoids high dimensionality.
- **Feature selection:** Mutual information is used because it is a filter-based method that can work with both classification and regression tasks. It helps keep the most informative features and remove weak features before model training.

This approach was chosen because it gives a good balance between automation, simplicity, and speed. More advanced techniques, such as K-nearest neighbour imputation, autoencoders, and wrapper feature selection, can be useful in specific cases, but they need more computation.

One of the important additions in this project is that the module is not only used for data preparation. It also supports exploratory data analysis report.

Also the module was extended to support text and image data preparation, not only tabular data.

2.3.2 Reporting in Pipeline

In this Part of the module, I implemented a simple reporting approach to help the user understand the pipeline results. Sometimes the pipeline produces many outputs, such as graphs, and different branch results. If these results are only shown in the terminal, they may be difficult to read and compare.

The general framework of the approach is as follows:

1. Take the results produced by the pipeline.
2. Take the graph from XML file.
3. Collect the metric results for each branch.
4. Put all the information in an organized HTML report.

The main idea of this module is to make the pipeline output easier to read. The report shows the important results in one place instead of leaving the user to search through terminal outputs or many files.

The main reporting parts are:

- **Pipeline graph:** The report first shows the full pipeline graph. This graph helps the user understand the general structure of the pipeline and how the branches are connected.
- **Branch nodes:** After that, the report shows the nodes inside each branch. This helps the user understand the steps that were applied in every branch separately.
- **Branch comparison:** By showing the metric results for all branches together, the report helps the user compare the branches and see which one performed better.

This approach was chosen because it makes the results clearer and easier to review. It also separates the reporting part from the main pipeline logic, which makes the system more organized.

2.3.3 Benchmarking

In this module, I implemented a benchmarking approach to compare machine learning submissions in a fair way. The module allows users to submit their prediction files and evaluates all submissions using the same test data, target values, and metric. This makes the comparison fair because all users are tested under the same conditions. The module also supports learning between users. After submissions are evaluated, users can view other users' code and results. This helps them understand different solution ideas, learn from stronger submissions, and improve their own work in future trials.

The general framework of the approach is as follows:

1. Create a benchmark task.
2. Separate the test data from the correct target values.
3. Allow users to submit their prediction files.
4. Check the submitted file format.

5. Compare the submitted predictions with the correct target values.
6. Calculate the score using the selected metric.
7. Show the results in a leaderboard.

The main idea of this module is to keep the evaluation fair. All users are tested using the same test data, the same target file, and the same metric.

The main benchmarking parts are:

- **Hidden target values:** The correct answers are kept hidden from the users to make the comparison fair.
- **Prediction submission:** Users only submit their predicted values.
- **Metric evaluation:** The system calculates the score using the selected evaluation metric.
- **Leaderboard:** The results are displayed in a leaderboard so users can compare their scores easily.

This approach was chosen because it gives a simple and fair way to compare different machine learning solutions. It also reduces manual checking and makes the evaluation process automatic.

Chapter 3: System Design and Architecture

This chapter explains the design of the modules included in this individual report. The main modules are Auto Preprocessing, Accelera Pipeline Reporting, and the Benchmarking Platform.

3.1 Overview and Assumptions

The design of these modules depends on some assumptions.

3.1.1 Auto Preprocessing Assumptions

For the Auto Preprocessing module, the assumptions depend on the type of dataset being processed. The module supports three types of data: tabular, text, and image datasets.

3.1.1.1 Tabular Dataset Assumptions

For tabular datasets, the Auto Preprocessing module assumes that:

- The dataset consists of numerical and categorical features, and the task is either classification or regression. Because we try to implement the standard structured machine learning problems where supervised learning is applied..
- The user provides configuration parameters to control preprocessing operations. This allows the pipeline to adapt to different datasets and user requirements.

3.1.1.2 Text Dataset Assumptions

For text datasets, the Auto Preprocessing module assumes that:

- we assume that the dataset has one text column and one categorical target label. This makes the pipeline easier to apply and keeps the number of features smaller after converting the text into numerical values.
- The user provides configuration parameters for text cleaning and feature extraction. To give the user a level of control over the preprocessing pipeline.

3.1.1.3 Image Dataset Assumptions

- The dataset is for image classification and follows this folder structure.

Training/

Cats/

Dogs/

Validation/

Cats/

Dogs/

- All images are stored in supported formats and are accessible. This ensures compatibility with image loading and preprocessing libraries.
- The user provides configuration parameters for preprocessing and augmentation. Because different datasets require different preprocessing strategies such as resizing, normalization, and augmentation.

3.1.2 Benchmark Assumptions

For the Benchmark module, the following assumptions are made. These assumptions are introduced due to system design constraints, particularly the limited server storage capacity for handling large dataset uploads.

- When the user creates a benchmark, they provide a tabular dataset. This ensures that the module operates on structured data suitable for standard machine learning evaluation.
- All datasets are stored in publicly accessible Google Drive links. This design choice is adopted to reduce server storage requirements.
- The user must provide three CSV files:
 - A training dataset containing both features and target labels.
 - A testing dataset containing the same feature structure as the training set but without the target column.
 - A ground truth target file containing only the target labels for the test set.

This separation ensures a standard evaluation protocol and prevents data leakage during benchmarking.

- When submitting predictions, the user provides a Google Drive link to a CSV file containing the predicted results. This ensures consistency in submission format across all users.
- The submitted prediction file must have the same length and format as the expected output to ensure correct evaluation and fair comparison between different users' models.

- When creating a benchmark, it is assumed that the user selects an evaluation metric that is appropriate for the problem type and provides a correct metric configuration. This is necessary to guarantee meaningful and valid performance measurement.
- For admin users, when defining an evaluation metric, it is assumed that a valid metric from the Scikit-learn library is selected. Because we use Scikit-learn metrics for evaluation.
- The selected metric must be properly configured and is expected to return a single scalar value for performance comparison and ranking purposes.
- The user is allowed to have only one active submission per benchmark. The user cannot create multiple submissions, but they are allowed to update their existing submission. This constraint is enforced to reduce storage usage.

3.1.3 Reporting Module Assumptions

The Reporting in Pipeline module depends on the following assumptions:

- The metric results are already generated from the pipeline modules.
- The metric result should be one of the supported types, such as a number, array, dictionary, string, or tuple.
- If the metric result needs to be displayed as a figure, the user should provide a plotting function.
- Figures are supported only in specific cases. If the metric result is a multi dimensional array or tuple and the user wants to display it as a graph, the user should provide a plotting function.
- The plotting function should take the metric result as input and return the correct Matplotlib `plt` object.
- The output folder path should be valid so the generated report and images can be saved.

3.2 System Architecture

This section explains the implemented modules in this report. It discusses the three main parts: Auto Preprocessing, Pipeline Reporting, and Benchmarking. For each module, the section explains module overview and how it works.

3.3 Auto Preprocessing Module

3.3.1 Module Overview

The module supports three dataset families:

- tabular datasets for classification and regression;
- text datasets for classification; and
- image datasets for classification.

The workflow starts from the raw data and user configuration. The module then applies suitable preprocessing steps and saves both the processed output and the fitted objects required for later use. It also has its own preprocessing reports. These reports are different from the pipeline report because they focus on the dataset itself: data overview, missing values, duplicate rows, column types, graphs, preprocessing decisions, selected features, and examples of the transformed data.

3.3.2 Tabular Training Flow




Figure 3.1: Tabular preprocessing workflow

The tabular training flow is implemented mainly in `classical_training_preprocessing.py`. The class receives a pandas DataFrame, target column, problem type, folder path, validation size, random state, thresholds, and report flag.

The training flow starts by validating the configuration. It checks the problem type, target column, validation size, random state, maximum ordinal threshold, cardinality threshold, missing value threshold, and feature importance threshold. It also saves the original input columns in `data_columns.pkl`, because the testing stage needs to compare new data against the same training structure.

The training flow is:

1. validate the input configuration
2. save original data columns for later testing
3. create a data overview for the report
4. remove rows has null target
5. remove duplicate rows
6. split data into training and validation sets
7. detect columns that should be dropped
8. detect each column type
9. generate graphs if reports are enabled
10. build a column transformer for feature preprocessing
11. preprocess the target according to classification or regression
12. apply feature selection using mutual information
13. save all fitted objects and metadata
14. generate the HTML report.

Algorithm 3.1: General tabular preprocessing flow

Input: DataFrame D, target column y, problem type, other configurations

Output: X_train, X_val, y_train, y_val

- 1: Validate configuration
 - 2: Store original dataset columns
 - 3: Build data overview and report information
 - 4: remove rows has null target
 - 5: Remove duplicate rows
 - 6: Split into training and validation sets
 - 7: Drop user-selected, high-missing, constant, and unsupported columns
 - 8: Detect column types from training data
 - 9: Generate feature and target graphs when report is enabled
 - 10: Fit feature preprocessing objects on X_train
-

```

11: Transform X_train and X_val
12: Fit target preprocessing object on y_train
13: Transform y_train and y_val
14: Select useful features using mutual information
15: Save fitted objects and metadata
16: Generate report if enabled

```

3.3.3 Data Overview and Cleaning

Before applying feature transformations, the module creates a detailed data overview. The original first rows are stored, then all string values are converted to lowercase and another preview is stored. This helps categorical and text like values become more consistent before splitting and encoding.

The overview also records dataframe information, numerical statistics, categorical statistics, missing values per column, duplicate count, duplicate percentage, and dataset shape. After this, rows with missing target values are removed, because the model cannot learn correctly without a target label. Then duplicate rows are dropped and the new shape is stored in the report data.

For classification tasks, the split is stratified when the target has more than one class and contains no null values. The goal is to keep the class distribution closer between training and validation data. For other cases, the normal `train_test_split` function is used with the configured `val_size` and `random_state`.

3.3.4 Tabular Column Type Detection

The module assigns a type to each feature and then chooses preprocessing based on the type. This rule based design is easy to understand and works without writing special code for every dataset.

For every column, the module calculates the data type, number of unique values, unique value ratio, missing value count, missing value ratio, constant value status, mode, and for numerical columns the minimum, maximum, median, and mean. A feature can be dropped if the user listed it in `columns_need_to_drop`, if it is constant, or if its missing value ratio is greater than `missing_threshold`.

Table 3.1: Detected tabular column types and preprocessing

Column type	Detection rule	Applied preprocessing
Binary	columns with exactly two unique values.	Most frequent imputation followed by ordinal encoding with unknown values mapped to -1.
Ordinal	Integer columns with unique values less than or equal to <code>max_unique_ordinal</code> .	Most frequent imputation followed by robust scaling.
Numerical	Integer columns with unique values greater than <code>max_unique_ordinal</code> .	Median imputation followed by robust scaling.
Continuous	Floating point columns.	Median imputation followed by robust scaling.
Low cardinality categorical	Object columns with unique values less than or equal to <code>cardinality_threshold</code> .	Most frequent imputation followed by binary encoding.
High cardinality categorical	Object columns with unique values greater than <code>cardinality_threshold</code> .	Most frequent imputation followed by frequency encoding.
Unsupported or other	Any unsupported data type.	Removed from the preprocessing pipeline.

3.3.5 Tabular EDA Graphs

The EDA part is generated per column, not as one general graph for the whole dataset. In `make_graphs`, the module loops over the training feature columns after their types are detected. For each column, it chooses the correct graph wrapper based on both the column type and the problem type. After all feature graphs are generated, it adds a target graph and a numeric correlation matrix.

Table 3.2: EDA graphs generated for tabular preprocessing

Column type	Problem type	Generated graphs
Binary and categorical	Classification	Null vs non null pie chart, category count plot, and category count split by target class.
Binary and categorical	Regression	Null vs non null pie chart, category count plot, and target box plot for each category.
Ordinal	Classification	Null vs non null pie chart, ordered count plot, and ordered count plot split by target class.
Ordinal	Regression	Null vs non null pie chart, ordered count plot, and target box plot for each ordinal value.
Numerical and continuous	Classification	Null vs non null pie chart, histogram with KDE, and feature box plot grouped by class.
Numerical and continuous	Regression	Null vs non null pie chart, histogram with KDE, and scatter plot between feature and target.
Target column	Classification	Null vs non null pie chart and target class count plot.
Target column	Regression	Null vs non null pie chart, target histogram with KDE, and target box plot.
Numeric columns	Both	Correlation heatmap.

For categorical columns with many values, the graph wrappers keep the top five categories and group the remaining values as `Other`. This keeps the graphs readable and avoids a very crowded x axis.

3.3.6 Target Processing

Target null rows are removed before the data split. After feature preprocessing, the target is processed according to the problem type. For classification, the target labels are encoded using `LabelEncoder`. For regression, the target values are scaled using `RobustScaler`. The fitted target preprocessor is saved as `target_preprocessor.pkl`, and target metadata is saved in `target_info.pkl`. In testing, this metadata is used to apply the same target transformation when labels are provided.

3.3.7 Feature Selection

After feature preprocessing, the module applies mutual information based feature selection. For classification it uses mutual information for classification tasks, and for regression it uses mutual information for regression tasks. The selected features are saved in `selected_features.pkl`. This reduces unnecessary feature columns and keeps the testing stage aligned with training.

3.3.8 Saved Tabular Artifacts

Table 3.3: Saved artifacts for tabular preprocessing

Artifact	Purpose
<code>data_columns.pkl</code>	Original input columns before preprocessing.
<code>col_drop.pkl</code>	Columns removed during training preprocessing.
<code>training_preprocessor.pkl</code>	Fitted feature transformer.
<code>feature_names.pkl</code>	Names of generated transformed features.
<code>target_preprocessor.pkl</code>	Fitted label encoder or target scaler.
<code>target_info.pkl</code>	Target metadata used during testing.
<code>bool_type_col.pkl</code>	Boolean columns that need consistent handling.
<code>selected_features.pkl</code>	Feature names selected after mutual information.

3.3.9 Tabular Testing Flow

The testing class loads the saved artifacts and applies the same transformations to new data. The testing data must not fit a new preprocessor. It should only use what was fitted during training. The test data is checked against the saved column list, dropped columns are removed, feature preprocessing is applied, selected features are kept, and the target is transformed if it exists.

3.3.10 Text Training Flow




Figure 3.2: Text preprocessing workflow

The text preprocessing class is used when the dataset contains a text column and a target column. It inherits shared tabular behavior for overview, duplicates, and splitting, but it uses text-specific feature preprocessing. During initialization, the class validates that the text column exists, that it is different from the target column, that the text column has object type, and that the target is suitable for classification. Any columns other than the text column and target column are removed because the text pipeline is designed for one text feature.

The text flow is:

1. generate a dataset overview
2. remove duplicate rows
3. split into training and validation sets
4. generate optional graphs such as target distribution and top word frequency
5. fill missing text values with an empty string

6. clean and tokenize text
7. transform text using TF-IDF
8. encode the target labels
9. save TF-IDF vectorizer, label encoder, and metadata
10. generate the report if enabled.

The report graphs for text are simpler than tabular column graphs. The `TextGraph` wrapper creates a null vs non null pie chart for the text column and a top word frequency plot. The target graph is generated using the classification target wrapper, so the user can also see the class distribution. This is useful because text datasets often have missing text, very frequent repeated words, and imbalanced labels.

3.3.11 Text Tokenization

The custom tokenizer lowercases the text, handles common contractions, keeps important negation words such as `not` and `no`, removes noise characters, separates digits from letters, removes short tokens, and applies Porter stemming. This prepares the text before TF-IDF vectorization. The fitted `TfidfVectorizer` is saved as `training_preprocessor.pkl`, while the target label encoder and metadata are saved for test-time reuse.

The tokenizer also normalizes special apostrophe characters and expands words such as `can't`, `won't`, and `shan't`. This detail matters because negation can change the meaning of a sentence. For example, removing the word `not` can make a negative review look positive, so the stopwords list keeps `not` and `no`.

Algorithm 3.2: Text tokenizer flow from the implementation

Input: raw text sentence

Output: cleaned token list

- 1: Convert text to lowercase and remove extra spaces
 - 2: Normalize apostrophe characters
 - 3: Expand common negative contractions
 - 4: Add spaces between letters and digits
 - 5: Remove characters except letters, digits, and apostrophes
 - 6: Split text into words
 - 7: Remove stopwords except "not" and "no"
 - 8: Remove very short tokens
 - 9: Apply Porter stemming
 - 10: Return cleaned tokens
-

After cleaning, the module fits TF-IDF on the training text only and transforms the validation text with the same vectorizer. This avoids fitting on validation data. The report stores a small before and after preview, where the original text is shown beside

the cleaned and stemmed tokens. For the target, the module uses `LabelEncoder` and saves it as `target_preprocessor.pkl`. It also saves `data_info.pkl` with the text column, and target column so the testing class can reuse the same configuration.

3.3.12 Text Testing Flow

The text testing class loads `data_info.pkl`, `training_preprocessor.pkl`, and `target_preprocessor.pkl`. It checks that the test dataframe contains the saved text column. If the target column is missing, it works in feature only mode and returns transformed features with `None` labels.

The testing flow lowercases the test data, fills missing text values with an empty string, and applies the saved TF-IDF vectorizer. If labels exist it transformed with the saved label encoder. This keeps the testing behavior close to training and prevents the test set from fitting its own vocabulary or label mapping.

3.3.13 Image Training Flow




Figure 3.3: Image preprocessing workflow

The image classification preprocessing class expects a folder structure where each class has its own folder. It scans the folders, builds class to label and label to class

mappings, validates images, optionally splits the training folder into training and validation sets, applies resizing and augmentation, and returns PyTorch data loaders.

The validation folder is optional. If it is provided, its class folders must match the training classes. If no validation folder is provided, the module can split the training folder using the configured validation size. Invalid or unreadable image files are not used for training, but they are recorded in the report.

The image path saves:

- class to label mapping
- label to class mapping
- image size information
- report data about valid and invalid images.

The image report includes label summary graphs, random sample images from the folders, and sample images after passing through the data loader. These graphs are important because they show whether resizing, and augmentation settings are working as expected.

The class mappings are created from the sorted training class names. This means the same class name always receives the same numeric label as long as the class folder names are the same. The mappings are saved in `class2label_mapping.pkl` and `label2class_mapping.pkl`. The selected image size is also saved in `data_info.pkl`. For each class folder, the module checks every image path using the image validation utility. Valid images are stored with their numeric labels. Invalid images are skipped, but their paths and labels are saved in the report data so the user can know which files were removed. If no valid image exists, the module raises an error instead of returning an empty loader.

The DataLoader construction is also different between training and validation. The training loader uses shuffling and can apply augmentation. The validation loader does not use augmentation and does not shuffle, because validation should measure model behavior on stable data. The available augmentation settings include horizontal flip, vertical flip, rotation, brightness change, and contrast change, controlled by the probability and parameter values passed to the constructor.

Algorithm 3.3: Image training preprocessing flow from the implementation

Input: training image folder, optional validation folder

Output: training and validation DataLoaders

- 1: Read class folders from the training directory
- 2: Create class to label and label to class mappings
- 3: Save mappings and image size information
- 4: Validate image files and skip invalid images
- 5: Build the data overview for the report

-
- 6: Split training images if no validation folder is provided
 - 7: Create training dataset with resize and optional augmentation
 - 8: Create validation dataset without augmentation
 - 9: Build DataLoaders
 - 10: Generate label, sample, and loader graphs
 - 11: Generate the image preprocessing report
-

3.3.14 Image Testing Flow

The image testing class is used for new images after the training preprocessing artifacts already exist. It loads the saved image size and class mapping, validates the provided image paths, and creates a test dataset and DataLoader. If class names are provided, each class name is converted to its saved numeric label. If class names are not provided, the loader can still be used for prediction only mode.

Testing does not apply augmentation. It only applies the same resizing and tensor conversion rules needed to make images compatible with the trained model. Invalid images are returned separately, which is useful because the user can see which inputs were ignored instead of silently losing files.

3.3.15 Preprocessing Reports

The Auto Preprocessing report belongs to the Auto Preprocessing module because it explains what happened to the dataset during preprocessing. It is separate from the Accelerera pipeline report, which explains pipeline branches and metric results.

For tabular data, `TabularPreprocessingReport` receives the `report_data` dictionary collected during preprocessing. The report contains:

- original and lowercased data previews;
- dataframe structure, numerical statistics, and categorical statistics;
- missing values, duplicate rows, and dataset shape;
- target-null handling and duplicate removal results;
- train-validation split sizes;
- dropped columns and the reason for each dropped column;
- EDA graphs generated for each feature column;
- preprocessing decision for every feature;
- processed samples of training and validation data; and
- feature-importance scores and selected features.

For text data, the same report class is reused in text mode. It shows the text overview, text graph, target graph, cleaning and tokenization preview, TF-IDF step, and encoded target samples. This helps show how raw sentences became a sparse numerical matrix.

For image data, `ImagePreprocessingReport` is used instead because the input is folder-based. It shows class names, class-to-label mapping, number of valid images, invalid image counts, invalid image paths, split information, label summary graphs, random sample images, and samples after `DataLoader` preprocessing.

This report is useful because preprocessing is usually a hidden part of a machine-learning project. With the report, the user can see the important changes instead of only receiving transformed arrays.

3.4 Accelera Pipeline Reporting Module

3.4.1 Module Overview

This chapter focuses only on the report classes related to `accelera_pipe`. These reports are different from the Auto Preprocessing report. The Auto Preprocessing report explains data cleaning and dataset transformation, while the pipeline report explains the executed pipeline graph, branches, selected path, and metric results.

The pipeline reporting part has three main layers. `ReportBase` provides the shared HTML page structure. `GraphPipelineReport` groups and displays pipeline metric results. `GraphReport` adds the graph image, branch tables, and best branch summary.

3.4.2 Accelera Report Base

`ReportBase` defines shared HTML report behavior. It stores the start and end HTML content, common styling, graph display helpers, and the `create_html_file` method that writes `report.html`. This class is important because different report classes can reuse the same structure instead of each one writing a full HTML page from the beginning.

The base report also has a helper for showing saved graph images. It receives a folder path and image names, checks whether the graph files exist, and adds the images to the HTML content. This keeps graph display consistent for any report that needs to include saved image files.

3.4.3 Pipeline Metric Report

`GraphPipelineReport` is a pipeline reporting class, not the same as the base Accelera report. It inherits from `ReportBase`, but its job is to format metric results produced by the pipeline execution. It receives a list of metric result dictionaries. Each result contains the metric name, the metric value, optional plotting function, labels, and headers.

The class groups results by metric name using `handle_metric`. After that, `metric_`

`display` chooses the suitable display wrapper based on the result type:

- scalar numbers are displayed with `DisplaySingleNumber`;
- one-dimensional arrays are displayed with `DisplayArraySingle`;
- multi-dimensional arrays are displayed with `DisplayMultiArray` or `DisplayFigure`;
- dictionaries are displayed with `DisplayDict`;
- strings are displayed with `DisplayString`; and
- tuples are displayed with `DisplayTupleNotCurve` or `DisplayFigure`.

3.4.4 Pipeline Graph Report

`GraphReport` is another pipeline report class. It extends `GraphPipelineReport`, but it adds graph-structure reporting. It reads the serialized XML graph file, extracts nodes and edges, renders a Graphviz image, finds branches through the graph, displays all branches, and also shows the best branch based on nodes marked as selected in the final path.

The generated graph image labels each node using node id, node type, and node name. Selected nodes are drawn in red, while the other nodes are drawn in blue. Metric nodes are also collected so that the metric summary can be matched with the graph branch where the metric was produced.

Algorithm 3.4: Pipeline graph report generation flow

Input: XML graph path, metric results

Output: HTML pipeline report and graph image

- 1: Parse the XML graph file
 - 2: Read nodes from XML layers
 - 3: Mark selected nodes and collect metric node ids
 - 4: Read edges and build graph connections
 - 5: Render the pipeline graph image with Graphviz
 - 6: Extract all branches until metric nodes are reached
 - 7: Group branches that share the same path
 - 8: Display all branches and the selected best branch
 - 9: Group metric results and format them by result type
 - 10: Write `report.html`
-

3.5 Benchmark Platform Module

3.5.1 Module Purpose

The Benchmarking module provides a shared place for evaluating machine-learning submissions. A benchmark creator defines the task, dataset links, target column, and

metric. Other users submit prediction files, and the system calculates the score. This keeps comparison more fair because all users are scored using the same target file and the same metric configuration.

The module is implemented in the `benchmark` folder. It has a Node.js and Express backend, MongoDB schemas using Mongoose, Python scripts for file validation and scoring, and a React frontend for benchmark and leaderboard screens.

3.5.2 Functional Description

When creating a benchmark, the user provides a title, description, problem type, target column, dataset link, test dataset link, ground truth target link, and evaluation metric. The module checks that the test dataset does not include the target column and that the target file includes the required target column. In the code, the target file is stored as `predictedColumnLink`, although it actually represents the true target labels used for scoring submissions.

When submitting to a benchmark, the user provides a prediction file and a repository link. The scoring process loads the target file and prediction file, checks that they have the same number of rows, and computes the selected metric. The submission is then saved and displayed in the leaderboard.

3.5.3 Backend Architecture

The backend server is started from `benchmark/backend/server.js`. It loads configuration, connects to MongoDB, enables CORS and JSON parsing, and mounts the main routes:

- `/metrics` for metric creation and listing;
- `/benchmark` for benchmark creation and browsing; and
- `/submission` for leaderboard submissions.

The backend uses schemas for benchmarks, metrics, and submissions. The route files handle benchmark browsing and creation, metric browsing and creation, and submission scoring. The Python scripts are called from the backend when the system needs to validate uploaded links or calculate a submission score.

3.5.4 Database Models

The main benchmark data is stored in `benchmark`, `metric`, and `submission` collections. The `metric` collection stores the display name, the scikit learn metric function name, needed parameters, problem type, and whether higher or lower values are better.

The `benchmark` collection stores the benchmark title, description, target column, dataset link, test set link, target label link, problem type, evaluation metric, metric parameters, creation date, and creator. The `submission` collection stores benchmark

id, submitter id, repository link, prediction file link, submission date, and score. It also has an index on benchmark id and score to support leaderboard access.

3.5.5 Main Backend Components

Table 3.4: Benchmarking module components

Component	Responsibility
Benchmark management	Create, list, filter, view, and delete benchmark tasks.
Metric management	Define metrics, problem type, parameters, and whether higher or lower score is better.
Submission management	Store prediction submissions, repository links, scores, submitter, and submission date.
Validation scripts	Check benchmark files and submission files before scoring.
Scoring scripts	Compute the selected scikit learn metric on predictions and ground truth.
Leaderboard interface	Display submissions ordered by the metric direction.

3.5.6 Benchmark Creation Flow

The benchmark creation endpoint is implemented in `routes/benchmark.js`. The backend first normalizes the problem type and checks that it is either classification or regression. It also checks that the benchmark title is unique. Dataset and file links are validated, and Google Drive file links are required for the test file and target file.

After link validation, the backend calls the Python validation script through `run_python`. The script downloads or retrieves the test and target files, then checks two important rules:

1. the test dataset must not contain the target column; and
2. the target file must contain the target column.

The current validation script also checks that the target column is numeric. If the validation succeeds, the benchmark is saved in MongoDB.

Algorithm 3.5: Benchmark creation flow

Input: benchmark form data

Output: saved benchmark or validation error

- 1: Receive benchmark creation request
- 2: Validate problem type
- 3: Check title uniqueness
- 4: Validate dataset, test set, and target label links

```

5: Run validate_user_links.py
6: If the test file contains the target column, reject benchmark
7: If the target file does not contain the target column, reject benchmark
8: Save benchmark document in MongoDB

```

3.5.7 Metric Management

Metrics are managed through `routes/metrics.js`. A metric is created by providing a display name, scikit learn metric function name, problem type, needed parameters, and `whichBetter`. The `whichBetter` value must be either `higher` or `lower`. This value is important because it controls leaderboard sorting.

The system prevents duplicate metrics with the same name, problem type, and scikit learn metric name. It also prevents deleting a metric if at least one benchmark still uses it.

3.5.8 Submission and Scoring Flow

The submission endpoint is implemented in `routes/submissions.js`. The module prevents the same submitter from creating more than one submission for the same benchmark. This is done by checking for an existing submission with the same benchmark id and submitter id. If the submitter wants to change the prediction file later, the update endpoint is used instead.

For each submission, the backend validates the repository link and prediction file link. Then it calls `get_score.py`. The Python script retrieves the true labels and submitted predictions, checks that the files have the same number of rows, checks that the prediction file contains the target column, and then calls the selected metric from `sklearn.metrics`.

Algorithm 3.6: Submission scoring flow

Input: benchmark target file, submitted prediction file, metric name
Output: score or validation error

```

1: Retrieve true target file
2: Retrieve submitted prediction file
3: Check that both files have the same number of rows
4: Check that submitted file contains the target column
5: Load metric function from sklearn.metrics
6: Compute metric(true_target, predicted_target, metric_parameters)
7: Return score to the backend as JSON
8: Save submission with score

```

The scoring algorithm is kept in Python because the project already uses Scikit learn metrics there. The backend sends the needed file links, target column, metric name, and metric parameters. The Python script returns either a score or an error message.

3.5.9 Leaderboard Ordering

Leaderboard ordering is handled by the submission route. When submissions are requested for a benchmark, the backend reads the benchmark metric and checks `whichBetter`. If the metric is better when higher, submissions are sorted by descending score. If the metric is better when lower, submissions are sorted by ascending score. This lets the same module support metrics such as accuracy and F1 score, where higher is better, and metrics such as mean absolute error, where lower is better.

3.5.10 Frontend Interface

The frontend is implemented with React under `benchmark/frontend/src`. The relevant pages for benchmarking are:

- `home.jsx`: entry page with links to benchmarks and metrics;
- `benchmarks.jsx`: lists benchmarks and filters them by user or problem type;
- `createBenchmark.jsx`: form for benchmark creation;
- `metrics.jsx` and `createMetric.jsx`: metric listing and metric creation;
- `displayBenchmark.jsx`: shows benchmark details; and
- `leaderBoard.jsx`: displays submissions and allows adding, updating, or deleting submissions when permitted.

3.5.11 Benchmark Constraints

The module has several important constraints:

- it currently supports classification and regression tasks;
- datasets and prediction files are provided as links;
- the test dataset must not contain the target column;
- the target file and prediction file must contain the selected target column;
- the target file and prediction file must have the same number of rows;
- metrics must map to valid scikit learn metric functions;
- benchmark creation and submission actions are protected by backend permission checks; and
- one user has only one active submission per benchmark, but the user can update it.

Chapter 4: System Testing and Verification

4.1 Testing Setup

The modules were tested using real datasets, saved preprocessing artifacts, generated reports, and benchmark workflow checks. The Auto Preprocessing results were evaluated by training machine learning models on the transformed outputs. Reporting was checked by generating HTML and graph outputs. Benchmarking was checked through benchmark creation, metric handling, submission validation, scoring, and leaderboard behavior.

4.2 Testing Plan and Strategy

The testing strategy focuses on verifying that each module works independently and that the generated outputs are usable by later steps. For preprocessing, the main check is that training artifacts are saved and can be reused during testing. For reporting, the main check is that report data is collected and rendered correctly. For benchmarking, the main check is that submissions are validated and scored using the same metric configuration.

4.3 Module Testing

4.3.0.1 Auto Preprocessing Module

Most of the preprocessing functions were tested using unit tests. These tests were used to check validation rules, artifact saving, and feature transformation. Real datasets were also used to test the system in a more practical way. These datasets were collected from Kaggle for tabular, text, and image data. For the tabular and text experiments, the data was split using a validation size of 0.2 and random state 42.

For tabular data, both classification and regression datasets were used. After preprocessing, the output was passed to a Random Forest model to check that the processed data was ready for training and validation. For text data, the processed TF-IDF output was passed to an `XGBClassifier`. This checked that the text preprocessing steps converted raw text into numerical features correctly. For image data, the output `DataLoaders` were passed to a ResNet model. This checked that the image preprocessing module created usable loaders.

Table 4.1: Datasets Used in Module Testing

Dataset Name	Shape	Dataset Type	Problem Type
startup_founder_burnout_2026	(50000, 29)	Tabular	Classification
healthy_diet_calorie_intake	(6000, 15)	Tabular	Classification
gaming_addiction	(250, 49)	Tabular	Classification
student_exam_performance_dataset	(10000, 17)	Tabular	Classification
heart	(1025, 14)	Tabular	Classification
Dataset_spine	(310, 7)	Tabular	Classification
diabetes_dataset	(100000, 31)	Tabular	Classification
student_placement_synthetic	(100000, 18)	Tabular	Classification
Titanic-Dataset	(891, 12)	Tabular	Classification
customer_purchase_data	(1500, 9)	Tabular	Classification
global_ai_jobs	(90000, 35)	Tabular	Regression
CarPrice_Assignment	(205, 26)	Tabular	Regression
Housing	(545, 13)	Tabular	Regression
PetImages	(24998,)	Image	Classification
DailyDialog	(11327, 2)	Text	Classification
TestReviews	(4321, 2)	Text	Classification

Table 4.2: Models and Evaluation Metrics Used in Module Testing

Dataset Name	Model Used	Evaluation Metric	Accelera Result
startup_founder_burnout_2026	RandomForest	Macro F1-score	0.9999
healthy_diet_calorie_intake	RandomForest	Macro F1-score	1.0000
gaming_addiction	RandomForest	Macro F1-score	1.0000
student_exam_performance_dataset	RandomForest	Macro F1-score	0.8573
heart	RandomForest	Macro F1-score	0.7860
Dataset_spine	RandomForest	Macro F1-score	0.7840
diabetes_dataset	RandomForest	Macro F1-score	0.9985
student_placement_synthetic	RandomForest	Macro F1-score	0.9986
Titanic-Dataset	RandomForest	Macro F1-score	0.7906
customer_purchase_data	RandomForest	Macro F1-score	0.9385
global_ai_jobs	RandomForest	R-squared Score	0.9227
CarPrice_Assignment	RandomForest	R-squared Score	0.9562
Housing	RandomForest	R-squared Score	0.6037
PetImages	ResNet model	Accuracy	0.9700
DailyDialog	XGBClassifier	Macro F1-score	0.9000
TestReviews	XGBClassifier	Macro F1-score	0.6080

These results show that the preprocessing output can be used by different models after transformation. The tabular results were strong in many datasets, especially the larger classification datasets and the car price regression dataset. The lower scores on datasets such as TestReviews show that preprocessing alone does not solve every modeling problem. Some datasets still need better features, model tuning, or more suitable models.

4.3.1 Comparison with AutoCleanML

The project also compares Accelera preprocessing with AutoCleanML on tabular datasets. The comparison uses several classification and regression datasets from Kaggle. For classification, the Macro F1 score is used. For regression, the R squared score is used. The same split settings are used, with validation size 20% and random state 42.

After preprocessing, the output data was tested using four models: Logistic Regression, Random Forest, Decision Tree, and KNN. The tables place the AutoCleanML and Accelera scores beside each other for the same model, so the comparison can be read directly.

Table 4.3: Classification comparison using Logistic Regression and Random Forest

Dataset	Logistic Regression		Random Forest	
	AutoCleanML	Accelera	AutoCleanML	Accelera
Startup founder burnout	1.0000	1.0000	0.9999	0.9999
Healthy diet calorie intake	0.9917	0.9890	1.0000	1.0000
Gaming addiction	0.8937	0.8937	0.9646	1.0000
Student exam performance	0.8667	0.8657	0.8579	0.8573
Heart	0.8019	0.8007	0.7689	0.7860
Spine dataset	0.8691	0.7943	0.8360	0.7840
Diabetes dataset	0.9995	0.9985	0.9996	0.9985
Student placement synthetic	0.5657	0.5879	0.6241	0.9986
Titanic dataset	0.7734	0.7864	0.7906	0.7906
Customer purchase data	0.8484	0.8375	0.9386	0.9385

Table 4.4: Classification comparison using Decision Tree and KNN

Dataset	Decision Tree		KNN	
	AutoCleanML	Accelera	AutoCleanML	Accelera
Startup founder burnout	1.0000	1.0000	0.9824	0.9881
Healthy diet calorie intake	1.0000	1.0000	0.9379	0.9184
Gaming addiction	1.0000	1.0000	0.6905	0.8512
Student exam performance	0.7987	0.8063	0.8423	0.8338
Heart	0.7213	0.7704	0.7689	0.7860
Spine dataset	0.7686	0.7632	0.8239	0.8042
Diabetes dataset	0.9995	0.9972	0.9925	0.9980
Student placement synthetic	0.5827	0.9968	0.5995	0.8348
Titanic dataset	0.6979	0.7877	0.7203	0.7562
Customer purchase data	0.8486	0.8449	0.8544	0.8806

Table 4.5: Regression comparison using Linear Regression and Random Forest

Dataset	Linear Regression		Random Forest	
	AutoCleanML	Accelera	AutoCleanML	Accelera
Global AI jobs	0.9022	0.7498	0.9288	0.9227
Car price assignment	0.8886	0.8943	0.9535	0.9562
Housing	0.6506	0.6482	0.6087	0.6037

From the classification and regression results, Accelera preprocessing achieved competitive performance compared to AutoCleanML. The comparison was tested us-

Table 4.6: Regression comparison using Decision Tree and KNN

Dataset	Decision Tree		KNN	
	AutoCleanML	Accelera	AutoCleanML	Accelera
Global AI jobs	0.8640	0.8412	0.7586	0.7510
Car price assignment	0.9269	0.8825	0.7905	0.7478
Housing	0.3819	0.3634	0.5312	0.5336

ing different types of models, including linear models, tree-based models, Random Forest, and KNN. In many cases, the results of Accelera were very close to AutoCleanML, and in some cases Accelera achieved better scores.

These results show that the implemented preprocessing steps in Accelera are able to prepare data in a valid way for different machine learning techniques. The performance was not limited to one specific model type, but appeared across several models and datasets. This means that Accelera successfully reached results close to an existing automated preprocessing tool, while still using its own preprocessing approach.

However, the results also show that no preprocessing tool gives the best score in all cases. The final performance can still depend on the dataset, the problem type, and the model used after preprocessing.

4.3.2 Pipeline Reporting Module Testing

The graph report is exercised through examples that build branch-based workflows, serialize the graph to XML, and run `GraphReport`. The expected result is a `report.html` file and a graph image showing the serialized graph structure. The report also checks that metrics are displayed according to their result type.

4.3.3 Benchmarking Module Testing

The Benchmarking module was tested manually by submitting different model results and checking whether the submitted results were saved and displayed in the leaderboard. The main checked behavior was that the module receives a submission, stores the dataset name, metric name, and score, and orders the leaderboard correctly.

4.3.4 Testing Datasets

Table 4.7: Examples of datasets used in Auto Preprocessing testing

Dataset	Data type	Problem type
startup_founder_burnout_2026	Tabular	Classification
healthy_diet_calorie_intake	Tabular	Classification
ai_student	Tabular	Classification
Titanic-Dataset	Tabular	Classification
global_ai_jobs	Tabular	Regression
CarPrice_Assignment	Tabular	Regression
Housing	Tabular	Regression
PetImages	Image	Classification
DailyDialog	Text	Classification
TestReviews	Text	Classification

References

Bibliography

- [1] Matthieu Komorowski, Dominic C. Marshall, Justin D. Saliccioli, and Yves Crutain. “Exploratory data analysis.” In *Secondary Analysis of Electronic Health Records*, pages 185–203. Springer, 2016.
- [2] Mohammad Hossin and Md Nasir Sulaiman. “A review on evaluation metrics for data classification evaluations.” *International Journal of Data Mining and Knowledge Management Process*, 5(2):1, 2015.
- [3] Navin Sabharwal and Amit Agrawal. “Introduction to Natural Language Processing.” In *Hands-on Question Answering Systems with BERT: Applications in Neural Networks and Natural Language Processing*, pages 1–14. Springer, 2021.

Chapter 5: Conclusions and Future Work

5.1 Faced Challenges

The main challenges were keeping preprocessing consistent between training and testing, supporting different dataset types inside one module, generating useful reports without mixing report logic with preprocessing logic, and designing a benchmark workflow that evaluates all submissions fairly.

5.2 Gained Experience

This work provided practical experience in automated data preprocessing, report generation, feature engineering, metric handling, backend workflow design, and testing machine learning support modules using real datasets.

5.3 Conclusions

This report documented the Auto Preprocessing module, the Reporting module, and the Benchmarking module in Accelerera. The Auto Preprocessing module automates repeated preparation steps for tabular, text, and image data and saves artifacts for consistent testing. The Reporting module uses HTML pages, tables, saved figures, and metric display wrappers to explain outputs. The Benchmarking module provides a shared evaluation workflow for comparing prediction submissions using the same metric and ground truth.

Together, these parts support a more complete data-science workflow. The user can prepare data, inspect generated reports, and compare final predictions in a benchmark. The implementation still has limitations, but it gives a clear base for future development and evaluation.

5.4 Future Work

5.4.1 Design Advantages

The main advantage of Auto Preprocessing is consistency. The same fitted objects used during training are saved and loaded during testing. This reduces the chance of data leakage and mismatched transformations. Another advantage is that the module supports more than one data type, so the same project can prepare tabular, text, and image datasets.

The main advantage of the Reporting module is explainability. The user can see data summaries, generated visualizations, preprocessing decisions, graph structure when available, and metric results in one place.

The Benchmarking module helps make comparison fairer because all submissions are evaluated using the same ground truth and metric configuration.

5.4.2 Current Limitations

There are also limitations:

- the preprocessing rules are fixed heuristics and may not be optimal for every dataset;
- the tabular preprocessing path assumes ordinary structured data and may need extensions for time-series or special domain data;
- text preprocessing uses classical TF-IDF, not transformer-based language models;
- image preprocessing prepares data loaders but does not automatically choose the best neural network architecture;
- graph report generation depends on graph serialization before report execution;
- benchmark files are provided through links, so file availability depends on external storage; and
- the benchmark testing described in the current report is mainly manual.

5.4.3 Possible Improvements

Future improvements can include adaptive preprocessing based on the model type, better handling of imbalanced datasets, more advanced text preprocessing, server-side benchmark file storage, stronger automated tests for the Benchmarking module, and more polished report templates. The graph report can also be extended to include execution time and memory information for each branch.

Chapter A: User Guide

This chapter gives simple examples for calling the main classes from the code. The examples are small, but they show the normal constructor calls and the main method used to start preprocessing or report generation.

A.1 Using Tabular Auto Preprocessing

For tabular data, the user loads the dataset as a Pandas DataFrame, chooses the target column, chooses the problem type, and gives an output folder. The output folder stores the fitted preprocessing objects and the report.

Algorithm A.1: Calling tabular training preprocessing

```
import pandas as pd
from accelera.src.auto_preprocessing.core.classical_training_preprocessing import (
    ClassicalTrainingPreprocessing,
)

df = pd.read_csv("dataset.csv")

preprocessor = ClassicalTrainingPreprocessing(
    df=df,
    target_col="target",
    problem_type="classification",
    folder_path="outputs/tabular_run",
    val_size=0.2,
    random_state=42,
    is_report=True,
)

X_train, y_train, X_val, y_val = preprocessor.common_preprocessing()
```

For regression, the main change is the problem type and target column.

Algorithm A.2: Calling tabular regression preprocessing

```
preprocessor = ClassicalTrainingPreprocessing(
    df=df,
    target_col="price",
    problem_type="regression",
    folder_path="outputs/regression_run",
)
```

For test data, the user should pass the same folder path. This makes the test data use the saved training transformers instead of fitting new transformers.

Algorithm A.3: Calling tabular testing preprocessing


```
import pandas as pd
from accelera.src.auto_preprocessing.core.classical_testing_preprocessing import (
    ClassicalTestingPreprocessing,
)

test_df = pd.read_csv("test.csv")

test_preprocessor = ClassicalTestingPreprocessing(
    df=test_df,
    folder_path="outputs/tabular_run",
)

X_test, y_test = test_preprocessor.common_preprocessing()
```

A.2 Using Text Auto Preprocessing

For text data, the user provides the text column and target column. The module cleans the text, applies TF-IDF, encodes the labels, and saves the needed objects.

Algorithm A.4: Calling text training preprocessing

```
import pandas as pd
from accelera.src.auto_preprocessing.core.text_training_preprocessing import (
    TextTrainingPreprocessing,
)

df = pd.read_csv("reviews.csv")

text_preprocessor = TextTrainingPreprocessing(
    df=df,
    target_col="label",
    text_col="review",
    folder_path="outputs/text_run",
    tfidf_max_features=500,
    is_report=True,
)

X_train, y_train, X_val, y_val = text_preprocessor.common_preprocessing()
```

For testing text data, the user again passes the same folder path. If the test data does not contain the target column, the class returns the transformed text features and None for labels.

Algorithm A.5: Calling text testing preprocessing

```
from accelera.src.auto_preprocessing.core.text_testing_preprocessing import (
    TextTestingPreprocessing,
)

test_text = pd.read_csv("reviews_test.csv")
```

```

text_test_preprocessor = TextTestingPreprocessing(
    df=test_text,
    folder_path="outputs/text_run",
)

X_test, y_test = text_test_preprocessor.common_preprocessing()

```

A.3 Using Image Auto Preprocessing

For image preprocessing, the dataset should be arranged as class folders. For example:

```

Training/
  Cats/
  Dogs/

```

The user passes the training folder and output folder. If there is no separate validation folder, the module can split the training folder.

Algorithm A.6: Calling image training preprocessing

```

from accelera.src.auto_preprocessing.core.classification_image_training_preprocessing
    import (
        ClassificationImageTrainingPreprocessing,
    )

image_preprocessor = ClassificationImageTrainingPreprocessing(
    training_folder_images="data/PetImages/Training",
    folder_path="outputs/image_run",
    split_training=True,
    val_size=0.2,
    images_size=(224, 224),
    batch_size=16,
    augment=True,
)

train_loader, val_loader = image_preprocessor.common_preprocessing()

```

For image testing, the user provides a list of image paths. Class names are optional, but if they are provided, they must match the saved training classes.

Algorithm A.7: Calling image testing preprocessing

```

from accelera.src.auto_preprocessing.core.classification_image_testing_preprocessing
    import (
        ClassificationImageTestingPreprocessing,
    )

image_paths = ["test/cat1.jpg", "test/dog1.jpg"]

```

```
class_names = ["Cats", "Dogs"]

image_test_preprocessor = ClassificationImageTestingPreprocessing(
    image_paths=image_paths,
    image_class_names=class_names,
    folder_path="outputs/image_run",
    batch_size=4,
)

test_loader, invalid_images = image_test_preprocessor.common_preprocessing()
```

A.4 Reading Auto Preprocessing Reports

After running preprocessing with `is_report=True`, the user opens the generated `report.html` file from the output folder. The report shows the dataset overview, missing values, duplicate rows, dropped columns, graphs, preprocessing decisions, selected features, and processed data samples.

A.5 Using Pipeline Reports

To generate a pipeline report, the pipeline graph should be serialized first. Then the user passes the XML graph path, output folder, and metric results to `GraphReport`.

Algorithm A.8: Calling the pipeline graph report

```
from accelera.src.accelera_pipe.wrappers.graph_report import GraphReport

report = GraphReport(
    folderpath="outputs/pipeline_report",
    xmlpath="outputs/pipeline_graph.xml",
    results=metric_results,
)

report.execute()
```

The output is a `report.html` file. It contains the graph image, branch tables, selected branch, and metric summary.

A.6 Using the Benchmark Platform

The benchmark platform is used from the web interface. The benchmark creator enters the title, problem type, target column, dataset link, test-set link, target-label link, and evaluation metric. A user submits a prediction file link and repository link. The backend calls the Python scoring script, saves the score, and displays the result in the leaderboard.

The prediction file must contain the target column expected by the benchmark.

Also, the number of rows in the prediction file must match the number of rows in the target-label file.