



Research Description of the Accelera Parallelizer and Pipeline Modules

Mohamed Ashraf Nesma Osama Mazen Adel Mohamed Khater
ID: 9220658 ID: 9220912 ID: 9220625 ID: 9220713

Supervisor: Dr. Salma Abdel monem

April 17, 2026

Abstract

Accelera is a hybrid Python/C++ machine learning framework designed to combine Python's modelling productivity with C++ execution infrastructure. The project contains graph-based machine learning pipelines, parallel branch execution, HTML reporting, AutoML preprocessing utilities, dataset retrieval, deployment support, benchmarking infrastructure, and a Linux-only C/C++ loop parallelizer. This report is organized around five project modules: `accelera_pipe`, the code parallelizer, AutoML, deployment, and the benchmark platform. The `accelera_pipe` module constructs and executes directed machine learning workflows using a Python API over a C++ graph backend. The code parallelizer analyzes source loops, extracts features, classifies loops, and injects OpenMP pragmas. The remaining module sections are reserved for future expansion.

Contents

1	Project Introduction	3
2	Accelera Pipe Module	3
2.1	Module Location and Responsibility	3
2.2	Graph Backend	3
2.3	Pipeline Builder API	4
2.4	Execution Semantics	4
2.5	Preprocessing and Model Fitting	4
2.6	Prediction Nodes	4
2.7	Metric Nodes	4
2.8	Branching and Merging	5
2.9	Executed Graphs	5
2.10	Reporting	5
2.11	Empirical Comparison with sklearn	6
2.12	Implemented Pipeline Optimizations	7
2.13	How to Use the Accelera Pipe Module	9

3	Code Parallelizer Module	10
3.1	Module Location and Responsibility	10
3.2	Native Loop Extraction	11
3.3	Feature Extraction	11
3.4	Vectorization for the Classifier	11
3.5	OpenMP Classifier Model	12
3.6	Loop Class Resolution	12
3.7	OpenMP Pragma Generation	13
3.8	Generator Model Trial and Rule-Based Design Rationale	13
3.9	Python-to-C++ Converter	14
3.10	Limitations	15
3.11	How to Use the Parallelizer	15
3.12	Hard-File Parallelizer Evaluation	16
3.13	Integration with Accelerate Pipe	16
4	AutoML Module	18
5	Deployment Module	18
6	Benchmark Platform Module	18
7	Module-Level View	18
8	Conclusion	19

1 Project Introduction

Accelera targets experimental and production-oriented machine learning workflows where a user wants to express a pipeline in Python while delegating graph execution and branch scheduling to a compiled backend. Its main design idea is to represent preprocessing, model training, prediction, merging, and metric computation as graph nodes. Independent branches can then be evaluated in parallel, making it convenient to compare alternative preprocessors, estimators, and metrics in a single run.

The repository also includes a code parallelizer that operates outside the ML pipeline. That subsystem parses C/C++ source code through Clang AST bindings, extracts loop fragments, computes structural and dependency-related features, consults a remote classifier service, and inserts OpenMP `parallel` for pragmas where the loop appears safe or useful to parallelize. If the input file is Python, a small AST-based converter first emits equivalent C++ for a limited subset of Python syntax before loop analysis begins.

2 Accelera Pipe Module

2.1 Module Location and Responsibility

The `accelera_pipe` module is implemented under `accelera/src/accelera_pipe/`. It provides the Python API for building graph pipelines and report objects. Its core files are:

- `core/pipeline_base.py`: wraps the C++ `graph.Graph` object and exposes shared configuration methods.
- `core/pipeline.py`: public builder API for preprocessing, model, prediction, metric, merge, and branch nodes.
- `core/executed_graph.py`: reusable executed graph wrapper.
- `core/metric.py`: abstract metric interface.
- `wrappers/supervised_metric.py` and `wrappers/unsupervised_metric.py`: concrete metric adapters.
- `core/report_base.py` and wrapper display classes: HTML report generation for pipeline graphs and metrics.

2.2 Graph Backend

`PipelineBase` imports the native `graph` module. If this module is missing, it raises an import error that instructs the user to build the C++ bindings and make them available in `PYTHONPATH`. When initialized, the base class creates or reuses a `graph.Graph` instance and enables parallel execution by default:

```
self.__graph = _graph if _graph is not None else graph.Graph()
self.__graph.enableParallelExecution(True)
```

The node type mapping bridges Python method names to C++ enum values: `PREPROCESS`, `MODEL`, `PREDICT`, `METRIC`, and `MERGE`.

2.3 Pipeline Builder API

`Pipeline` is the main user-facing class. Each method either appends a node to the graph or, when `branch=True`, returns a lightweight `Node` object that can be supplied to `branch()`.

Table 1: Main `Pipeline` methods.

Method	Purpose	Important parameters
<code>preprocess</code>	Add a transformation step.	<code>name</code> , <code>func</code> , <code>branch</code> .
<code>model</code>	Add a trainable model.	<code>name</code> , <code>model</code> , <code>branch</code> .
<code>predict</code>	Add prediction over test data.	<code>test_data</code> , <code>output_func</code> , <code>positive_class</code> .
<code>metric</code>	Add a supervised or unsupervised metric.	<code>metric_name</code> , <code>y_true</code> , <code>plot_func</code> , <code>labels</code> , <code>headers</code> , extra metric parameters.
<code>merge</code>	Combine branch predictions.	<code>strategy</code> , currently defaulting to <code>hard_voting</code> .
<code>branch</code>	Split the graph into alternative paths.	A branch name and one or more <code>Node</code> objects created with <code>branch=True</code> .

2.4 Execution Semantics

Calling a `Pipeline` executes the underlying graph:

```
predictions, executed_graph = pipe(X, y, select_strategy="all")
```

The C++ backend returns an executed graph plus result values. The Python layer wraps the first returned item in `ExecutedGraph`; all remaining returned items are predictions or metric outputs. The selection strategy can be `all`, `max`, `min`, or a user-provided `custom_strategy`. Tests in the module show that these strategies are used to select paths according to metric outputs.

2.5 Preprocessing and Model Fitting

`Preprocess` nodes store the transformation object or function together with `execute_fit`, a helper that determines whether a component should be fit with `X` only or with both `X` and `y`. This lets the pipeline handle unsupervised transformers such as scalers and supervised transformers such as feature selectors. `Model` nodes similarly store the model object and fitting helper, allowing `sklearn`-compatible estimators and custom `Accelerate` estimators to participate in the graph.

2.6 Prediction Nodes

`predict(name, test_data, output_func="predict", positive_class=-1)` records the data on which the trained model should be evaluated. The `output_func` parameter can request outputs such as `predict` or `predict_proba`. When probabilities are used, `positive_class` can select one probability column.

2.7 Metric Nodes

The pipeline maps metric names to functions from `sklearn.metrics`. Then the metric-class helper inspects the metric function signature:

- A metric with true labels and predicted labels/scores becomes a **SupervisedMetric**.
- A metric with **X** and **labels** becomes an **UnSupervisedMetric**.
- Any incompatible signature is rejected.

SupervisedMetric.execute validates that **y_pred** and **y_true** have the same number of samples, then calls the sklearn metric. **UnSupervisedMetric.execute** validates predictions against the stored input **X**. Both return a dictionary containing the metric name, result, plot function, label names, and header names.

2.8 Branching and Merging

branch() accepts one or more branch node objects. Each branch can be a single node or an iterable of nodes. The method validates that all branch members are **Node** instances and rejects merge nodes inside branches because merge-in-branch behavior is not implemented yet. The validated branch objects, node types, and node names are passed to the C++ graph through **split()**.

Merge nodes combine multiple branch outputs. The default strategy is **hard_voting**. The test suite covers hard-voting merges, merges after preprocessing branches, probability prediction merges, single-model merges, and reuse of executed graphs after merging.

2.9 Executed Graphs

ExecutedGraph wraps a graph that has already been fitted or selected. Calling it with new **X** reuses the fitted pipeline. If **y_true** is provided, metrics are temporarily enabled for that call and disabled afterward.

Executed graphs can now be persisted as a single pickle file. Calling **save()** with no arguments creates **pipeline.pkl** in the current working directory. The user may also pass an explicit file path such as **"my_pipeline.pkl"**. Loading is performed through **ExecutedGraph.load()**, which reconstructs the native C++ graph and restores the fitted Python objects stored in the graph nodes. The loaded graph keeps the executed-state flag, so applying data to it performs inference directly rather than fitting the models again.

```
predictions, executed_graph = pipeline(X_train, y_train)

executed_graph.save()           # writes pipeline.pkl
loaded_graph = ExecutedGraph.load() # reads pipeline.pkl

test_predictions = loaded_graph(X_test)
```

The implementation uses a pybind pickle bridge for the C++ **Graph** object. During saving, the graph is converted into a Python dictionary containing node metadata, source-edge indices, selected-path flags, fitted preprocessors, and fitted model state. During loading, this dictionary is used to rebuild the graph topology and restore each node. This design stores the pipeline in one file, but it requires the Python objects inside the pipeline to be pickle-compatible; standard sklearn transformers and estimators satisfy this requirement, while lambdas and some local custom functions may not.

2.10 Reporting

The reporting layer creates HTML reports rather than PDF reports. The base class **ReportBase** defines the HTML shell and utility functions. **GraphReport** reads serialized pipeline XML, builds a Graphviz diagram, colors selected path nodes red and other nodes blue, lists branches, and then delegates metric rendering to **GraphPipelineReport**.

The graph-pipeline report groups metric records by metric name and chooses a display wrapper according to the result type:

- scalar values: `DisplaySingleNumber`;
- one-dimensional arrays: `DisplayArraySingle`;
- multi-dimensional arrays: `DisplayMultiArray` or `DisplayFigure`;
- dictionaries: `DisplayDict`;
- strings: `DisplayString`;
- tuples: `DisplayTupleNotCurve` or `DisplayFigure`.

`ModelReport` reuses this metric rendering and can optionally plot training history when a dictionary of history arrays is provided.

2.11 Empirical Comparison with sklearn

To evaluate the execution behavior of `accelera_pipe`, a controlled comparison was performed against two sklearn baselines. The first baseline uses a direct manual sklearn loop, where each candidate scaler and model is cloned, fitted, evaluated on validation data, and then the selected candidate is evaluated on test data. The second baseline uses `sklearn.pipeline` with the same candidate pairs. The `Accelerate` version expresses the same candidate space as two graph branches: two preprocessing alternatives (`StandardScaler` and `MinMaxScaler`) followed by two model alternatives (`LogisticRegression` and `SVC`). Therefore, all systems evaluate four candidate pipelines and select the path with the highest validation accuracy.

The experiment uses synthetic four-class classification data with 25 features. Each dataset is split into 60% training, 20% validation, and 20% test data. The validation split is used only for candidate selection, while the test split is used to report final accuracy. Table 2 shows that all three implementations selected an equivalent best candidate and obtained the same test accuracy in every run. For small inputs, the native graph overhead dominates; for larger inputs, the parallel branch execution of `accelera_pipe` reduces latency compared with both sklearn baselines. The memory measurements also show that `accelera_pipe` currently pays a larger memory cost because multiple graph branches and fitted node states can be alive at the same time. The largest run, with 150,000 training samples, 50,000 validation samples, and 50,000 test samples, is especially important: `accelera_pipe` matched the sklearn test accuracy while reducing runtime by approximately 309 seconds compared with both sklearn baselines.

Table 2: Fair comparison between manual sklearn, sklearn Pipeline, and `accelera_pipe`. Time is measured in seconds and RSS is the process resident-memory delta in MB.

Samples	Train	Val	Test	sklearn Time	sklearn RSS	sklearn Pipe Time	sklearn Pipe RSS	Accelerate Time	Accelerate RSS
1,000	600	200	200	0.06	1.62	0.05	0.16	0.34	19.32
5,000	3,000	1,000	1,000	2.47	3.62	2.12	1.09	2.02	21.62
10,000	6,000	2,000	2,000	2.64	11.25	2.67	2.19	2.10	28.43
50,000	30,000	10,000	10,000	23.96	126.95	23.25	9.75	14.81	209.25
100,000	60,000	20,000	20,000	77.79	139.70	77.10	47.65	67.04	300.93
250,000	150,000	50,000	50,000	803.56	303.95	803.01	28.52	494.64	594.86

Table 3: Accuracy and selected candidate behavior. The test-accuracy delta is computed as `accelera_pipe` minus manual sklearn.

Samples	sklearn Best Candidate	Accelera Selected Candidate	Validation Accuracy	Test Accuracy	Test Accuracy Delta
1,000	MinMaxScaler → SVC	MinMaxScaler → SVC	0.7900	0.8250	+0.0000
5,000	MinMaxScaler → SVC	MinMaxScaler → SVC	0.9160	0.9100	+0.0000
10,000	MinMaxScaler → SVC	MinMaxScaler → SVC	0.9385	0.9260	+0.0000
50,000	MinMaxScaler → SVC	MinMaxScaler → SVC	0.9565	0.9598	+0.0000
100,000	StandardScaler → SVC	StandardScaler → SVC	0.9675	0.9709	+0.0000
250,000	StandardScaler → SVC	StandardScaler → SVC	0.9774	0.9768	+0.0000




Figure 1: Latency scaling for the same four-candidate search space.




Figure 2: RSS memory delta for the same benchmark runs.

2.12 Implemented Pipeline Optimizations

Several optimizations were added to reduce unnecessary work and memory retention in the native graph backend. These changes target the main cost of branch-based pipeline search: multiple candidate paths can hold similar intermediate arrays, fitted states, and prediction outputs at the same time.

Optional cache execution. Caching for preprocess and model nodes was made opt-in. The public Python API keeps the parameter `cache=False` by default for both `preprocess()`

and `model()`. When caching is disabled, the C++ nodes skip the expensive `joblib.hash()`, `joblib.load()`, and `joblib.dump()` path. This reduces latency and memory pressure in normal search runs, while preserving cache support for repeated experiments where the user explicitly enables it.

Consumer-count based intermediate release. The graph backend now computes how many downstream nodes still consume each node output. After a node finishes, the consumer count of its source nodes is decremented. When the count reaches zero, temporary data for releasable nodes is cleared. This is especially important for preprocessing outputs, because a transformed training array may be needed by multiple model branches but becomes dead once the last model has finished fitting.

Input training-data release. The original input node stores the raw training X and y . After all downstream training consumers have executed, these references are cleared as well. This prevents the graph from retaining the original training arrays after fitted preprocessors and models no longer require them. The release is disabled for executed graphs during inference, because the executed graph input must remain available while prediction nodes apply the fitted preprocessing chain to new test data.

Separated output-container creation from input copying. Originally, the internal flag `should_create_new_data` controlled two different behaviors: creating a fresh branch output container and copying X/y before preprocessing. These meanings were separated into two concepts. A first preprocessing node in a branch still creates its own output dictionary, but the decision to defensively copy the input is controlled independently. This avoids corrupting branch storage semantics while allowing safe copy reduction.

Guarded defensive-copy reduction for preprocessing nodes. The graph separates branch-output creation from input-copy decisions through the `shouldCopyPreprocessInput()` helper. A return value of `true` means that the preprocessing node receives defensive copies of X and y ; a return value of `false` means that the node can safely read the existing branch input without copying it first. The helper first rejects non-preprocessing nodes, then inspects the Python preprocessing object stored under the `func` key. Non-sklearn callables remain conservative and are copied, because an arbitrary user function may mutate its input in place.

For sklearn transformer instances, copying is disabled only when the transformer is not configured for in-place mutation. The implementation calls `get_params(deep=False)` when available and checks the sklearn `copy` parameter. If `copy=False`, the graph keeps the defensive copy because the transformer is explicitly allowed to mutate its input. If the parameter is absent or inspection succeeds without finding an unsafe configuration, the copy is skipped. If inspection raises a Python error, the backend clears the error and falls back to copying. This design reduces memory for common sklearn preprocessors such as `StandardScaler` and `MinMaxScaler` under their default settings, while preserving isolation for custom callables and sklearn transformers configured with `copy=False`.

Executed-graph data retention reduction. Executed graph cloning keeps the fitted model state needed for future prediction, but avoids carrying unnecessary training intermediates such as raw training arrays and preprocessing output dictionaries. The selected executable path therefore keeps fitted preprocessors and fitted models as reusable state, while temporary training data can be released.

Duplicate first-branch-node sharing. The split operation now deduplicates identical first nodes created from the same split source. For example, if a branch is declared as:

```

pipe.branch(
    "preprocessing",
    pipe.preprocess("a", MinMaxScaler(), branch=True),
    pipe.preprocess("b", StandardScaler(), branch=True),
    pipe.preprocess("c", MinMaxScaler(), branch=True),
)

```

the graph creates one `MinMaxScaler` node and one `StandardScaler` node rather than two equivalent `MinMaxScaler` nodes. The signature used for this sharing depends on the node type and Python object representation, not on the user-facing node name, so identical operations with different names can still be shared.

The optimization is intentionally restricted to first branch nodes. For list branches, only the first item is eligible for sharing:




Figure 3: Branch graph before and after duplicate first-node sharing.

Later nodes are not merged, even if they have the same operation name or object representation, because their inputs may differ. For example, `[p1, p2, p3]` and `[p4, p2, p3]` keep separate `p2` nodes because one `p2` receives data from `p1`, while the other receives data from `p4`. This preserves correctness while removing obvious duplicate work at branch entry points.

Fair benchmark protocol. The comparison script was adjusted so manual `sklearn`, `sklearn Pipeline`, and `accelera_pipe` follow the same protocol: train and validate all four candidate paths, select the best validation path, and evaluate the test split only once for the selected path. This makes latency comparisons more meaningful by avoiding extra test predictions in the `sklearn` baselines.

2.13 How to Use the Accelera Pipe Module

A minimal classification pipeline is:

```

from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

from accelera.src.accelera_pipe.core.pipeline import Pipeline

X, y = make_classification(n_samples=1000, n_features=20, random_state=42)
X_test, y_test = X[:100], y[:100]

pipe = Pipeline()
pipe.preprocess("scale", StandardScaler())

```

```

pipe.model("logreg", LogisticRegression(max_iter=1000))
pipe.predict("predict", test_data=X_test)
pipe.metric("accuracy", "accuracy_score", y_true=y_test)

results, executed_graph = pipe(X, y)
print(results)

```

A branching experiment compares preprocessors and models:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import MinMaxScaler, StandardScaler

pipe = Pipeline()
pipe.branch(
    "preprocessing",
    pipe.preprocess("standard", StandardScaler(), branch=True),
    pipe.preprocess("minmax", MinMaxScaler(), branch=True),
)
pipe.branch(
    "models",
    pipe.model("logreg", LogisticRegression(max_iter=1000), branch=True),
    pipe.model("forest", RandomForestClassifier(n_estimators=50), branch=True),
)
pipe.predict("predict", test_data=X_test)
pipe.metric("accuracy", "accuracy_score", y_true=y_test)

results, executed_graph = pipe(X, y, select_strategy="max")
best_result = executed_graph(X_test, y_test)

```

To generate a graph report:

```

from accelera.src.accelera_pipe.wrappers.graph_report import GraphReport
from accelera.src.utils.accelera_utils import serialize

results, executed_graph = pipe(X, y, select_strategy="max")
serialize(pipe, "pipeline.xml")

report = GraphReport("pipeline_report", "pipeline.xml", results)
report.execute()

```

3 Code Parallelizer Module

3.1 Module Location and Responsibility

The parallelizer is implemented mainly in `accelera/src/utils/parallelizer.py`. Its Python-to-C++ helper is in `accelera/src/utils/py2cpp-converter.py`. Native loop extraction is provided by the `pybind` module `code_parallelizer_utils`, which is built from C++ sources under `src/ast/` and `src/utils/code_parallelizer_utils.cpp`. The CMake configuration enables these bindings only on Linux with LLVM/Clang available.

Conceptually, the module implements the following pipeline:

1. Read a C, C++, or supported Python file.
2. Convert Python input to C++ when the file extension is `.py`.
3. Extract loops using Clang-backed native bindings.

4. Cache extracted loop metadata in `.accelera_cache`.
5. Compute a 40-dimensional feature vector for each `for` loop.
6. Send the vector to the classifier endpoint configured by `ACCELERA_CLASSIFIER_ENDPOINT`.
7. Resolve the final loop class with local heuristics.
8. Generate OpenMP pragmas and replace the original loop text.
9. Write `parallelized_<name>.c` beside the input file and format it with `clang-format` when available.

3.2 Native Loop Extraction

The public loop extraction helpers wrap the native `code_parallelizer_utils` bindings. The default Clang argument is `-std=c++17`. The expected extracted loop records include at least the loop source code, loop type, start line, and end line. Although the C++ AST layer can identify several loop kinds, the Python insertion logic currently parallelizes only loops whose extracted type is `for`.

3.3 Feature Extraction

The function `extract_features(code)` performs lightweight static analysis using regular expressions after comment removal. The extracted features are intentionally pragmatic: they describe loop structure, reduction patterns, memory access shape, control flow, and rough computational intensity. Important feature groups include:

- **Reduction features:** detection of scalar `+=`, `x = x + ...`, `*=`, `x = x * ...`, and textual `min/max` usage.
- **Dependency features:** loop-carried dependency patterns such as `a[i - 1]`, read-after-write flags, write-after-read flags, and dependency distance.
- **Memory features:** array reads, array writes, indirect indexing, number of memory operations, and hashed array-name embeddings.
- **Control-flow features:** branch count, early exits via `break`, `continue`, or `return`, and function call counts.
- **Shape features:** constant loop bounds, estimated iteration count, consecutive nested loop depth, and nested brace depth.
- **Vectorization hint:** a boolean flag that becomes true when the loop has no branches, no function calls, no loop-carried dependency, and no indirect access.

3.4 Vectorization for the Classifier

`vectorize_features(features)` maps the feature dictionary to a 40-element `numpy.float32` vector. Boolean features occupy positions 0–16. Counts and normalized scalar descriptors occupy later positions. Array names are transformed into compact deterministic hash vectors rather than sent as raw strings. This representation is sent as JSON to the remote classifier:

```
response = requests.post(
    classifier_endpoint,
    json={"embedding": features.tolist()},
    timeout=config.REQUEST_TIMEOUT_S,
)
```

The classifier result is expected to contain `result`, which is then interpreted as a predicted loop class.

3.5 OpenMP Classifier Model

The classifier used by the Parallelizer is stored in `models/open_mp_classifier`. It is packaged as a small FastAPI service so that the Python parallelization logic can send a loop embedding and receive a loop category without embedding the model code directly inside the converter. The service loads two serialized files: `model.pkl`, which is the trained classifier, and `label_encoder.pkl`, which maps numeric model outputs back to textual labels.

The prediction API accepts one 40-dimensional feature vector:

```
{
  "embedding": [0.0, 1.0, 0.0, "..."]
}
```

and returns one of the supported loop labels:

```
{
  "result": "parallel_for"
}
```

The deployed model is an `XGBClassifier`. The saved model inspected from `model.pkl` has 300 boosting rounds, maximum tree depth 9, learning rate 0.03, objective `multi:softprob`, and 40 input features. The label encoder stores three classes: `none`, `parallel_for`, and `reduction`. The accompanying notebook `m1.ipynb` trains the model from `data/data.csv`, where each row contains 40 extracted static features and one label column. The dataset contains 32,595 loop samples: 15,240 labeled `none`, 13,482 labeled `parallel_for`, and 3,873 labeled `reduction`. Training uses a stratified 80/20 split with `random_state=42`.

Table 4: OpenMP classifier test metrics from the training notebook.

Class	Precision	Recall	F1-score	Support
<code>none</code>	0.82	0.85	0.84	3048
<code>parallel_for</code>	0.86	0.81	0.83	2696
<code>reduction</code>	0.79	0.82	0.80	775
Accuracy			0.83	6519
Macro average	0.82	0.83	0.83	6519
Weighted average	0.83	0.83	0.83	6519

These metrics show that the classifier is useful but not treated as the only source of truth. Accelera keeps deterministic fallback rules around the model: the classifier provides the first prediction, then rule-based checks can recover well-known safe cases such as reductions, collapsed nested loops, and independent array writes. This hybrid design is important because an incorrect OpenMP pragma can change program semantics, so the system favors explainable static evidence when the model is uncertain or unavailable.

3.6 Loop Class Resolution

The function `_resolve_loop_class(loop_code, pred_class)` combines remote prediction with local safety heuristics. If the classifier predicts a class other than `none`, that class is used. If it predicts `none`, the local logic can still upgrade the loop to:

- `reduction`, when scalar reduction updates are detected.

- `parallel_for`, when consecutive nested `for` loops are detected.
- `parallel_for`, when the loop performs independent array writes without loop-carried dependencies, indirect indexing, or early exits.

Otherwise the loop remains `none`. This means the module is not purely ML-based; it is a hybrid predictor with deterministic fallbacks for common OpenMP-friendly patterns.

3.7 OpenMP Pragma Generation

`_generate_omp_pragma_with_loop` creates the final annotation. For ordinary parallel loops it emits:

```
#pragma omp parallel for
for (...) {
    ...
}
```

For consecutive nested loops it adds `collapse(depth)`. For scalar addition or multiplication reductions it adds `reduction(+ : var)` or `reduction(* : var)`. The generator therefore currently focuses on OpenMP `parallel for`, `collapse`, and scalar reduction clauses.

3.8 Generator Model Trial and Rule-Based Design Rationale

An earlier generator trial attempted to learn OpenMP pragma generation as a sequence-to-sequence problem. In that formulation, the input sequence contained the loop source code prefixed by a class token such as `[CLS:parallel_for]` or `[CLS:reduction]`, and the target sequence was the expected OpenMP pragma. The implemented model was a custom PyTorch encoder-decoder network:

- **Encoder:** a bidirectional LSTM with embeddings, dropout, packed variable-length input handling, and layer normalization.
- **Decoder:** an LSTM decoder conditioned on Bahdanau additive attention over encoder states.
- **Bridge:** linear projections that map the bidirectional encoder hidden and cell states into the decoder state space.
- **Training:** teacher forcing, cross-entropy loss with padding ignored, Adam optimization, gradient clipping, and a ReduceLROnPlateau scheduler.

The deployed generator service in `models/open_mp_generator/app.py` uses the following configuration: embedding size 128, hidden size 256, three LSTM layers, dropout 0.2, greedy decoding, maximum input length 1500, and a checkpoint saved as `best_model.pth`. The tokenizer is a custom BPE tokenizer rather than a pretrained Hugging Face tokenizer. It is stored in `tokenizer.json`, uses the special tokens `<PAD>`, `<UNK>`, `<SOS>`, and `<EOS>`, and was configured for a target vocabulary of 8000 tokens. The saved tokenizer contains 8002 tokens and 7888 learned BPE merges.

Although this neural generator is expressive, two empirical limitations made a pure learned generator unsuitable for the current system. First, the available training data was small for a code-generation task and contained only a limited effective number of examples for some pragma classes. In particular, the corpus included about 15k negative examples and a comparatively narrow distribution of positive pragma patterns, which made the sequence model prone to memorization and overfitting. As a result, validation performance did not reliably imply generalization to unseen loop structures, variable names, dependency patterns, or clause combinations.

Second, expanding the corpus with AI-generated examples introduced distribution bias. Synthetic loops and pragmas tend to reflect the style, templates, and regularities of the generating tool. A model trained heavily on such data can learn those artifacts instead of learning robust OpenMP semantics. This lowers real-world generalization even when apparent training or validation error looks acceptable on similarly generated samples.

For these reasons, the implemented generator path was changed to a rule-based approach. OpenMP pragmas are now constructed from explicit static features such as consecutive loop depth and scalar reduction variables. This design is less flexible than open-ended neural generation, but it is more predictable, auditable, and aligned with the current objective: capture the most common and safe pragma forms, including `parallel for`, `collapse(n)`, and scalar `reduction` clauses, without hallucinating unsupported clauses or unbound variables.

3.9 Python-to-C++ Converter

The converter in `py2cpp_converter.py` uses Python’s built-in `ast` module. It does not attempt to be a complete Python compiler. Instead, it emits simple C++ for a restricted subset that is useful for loop experiments. Unsupported syntax raises `ValueError`. The output always contains an `int main()` for top-level statements, while Python function definitions are emitted before `main()` with `auto` return type and `auto` parameters.

Table 5: Supported operations and methods in the Python-to-C++ converter.

Category	Supported form	Generated C++ / Notes
Constants	<code>None</code> , <code>bool</code> , <code>int</code> , <code>float</code> , <code>string</code>	<code>nullptr</code> , <code>true/false</code> , numeric literals, escaped <code>std::string</code> literals.
Names	<code>x</code>	Emitted as the same identifier.
Binary operations	<code>+</code> , <code>-</code> , <code>*</code> , <code>/</code> , <code>%</code> , <code>**</code>	Parenthesized infix expressions. Power expressions are emitted as <code>std::pow(left, right)</code> with <code><cmath></code> .
Unary operations	unary <code>+</code> , unary <code>-</code> , <code>not</code>	Mapped to <code>+</code> , <code>-</code> , and <code>!</code> .
Comparisons	<code>==</code> , <code>!=</code> , <code><</code> , <code><=</code> , <code>></code> , <code>>=</code>	Only one comparison operator is allowed; chained comparisons are rejected.
Boolean operations	<code>and</code> , <code>or</code>	Mapped to <code>&&</code> and <code> </code> .
Function calls	<code>f(a, b)</code>	Emitted as ordinary C++ calls. Keyword arguments are generally unsupported.
Printing	<code>print(a, b)</code>	Mapped to <code>std::cout << a << " " << b << std::endl</code> . Print keywords are rejected.
Attributes	<code>obj.attr</code>	Mapped to <code>obj.attr</code> .
Subscript	<code>arr[i]</code>	Mapped to <code>arr[i]</code> . Slices are rejected.
Assignments	<code>x = expr</code>	First assignment declares inferred type; later assignments reuse the variable. Subscript assignment such as <code>X[i][j] = expr</code> is emitted directly. Multi-target assignments are rejected.

Category		Supported form	Generated C++ / Notes
Augmented assignment	assign-	<code>x += y, x -= y, x *= y, x /= y, x %= y</code>	Simple-name targets only. <code>x += 1</code> is emitted as <code>x++</code> .
Return		<code>return</code> or <code>return expr</code>	Emitted directly.
Conditionals		<code>if/else</code>	Standard C++ braces.
For loops		<code>for i in range(...)</code>	Supports <code>range(stop)</code> , <code>range(start, stop)</code> , and <code>range(start, stop, step)</code> with simple name targets.
Functions		<code>def f(a, b): ...</code>	Emitted as a C++17 template function with <code>auto</code> return type, for example <code>template <typename T0, typename T1> auto f(T0 a, T1 b) {...}</code> . Decorators are rejected.
Type inference		<code>bool, int, float, string</code> constants	Emits <code>bool</code> , <code>int</code> , <code>double</code> , or <code>std::string</code> ; otherwise <code>auto</code> .

3.10 Limitations

The parallelizer is Linux-oriented because it depends on LLVM/Clang development libraries and pybind bindings. It also depends on a remote classifier endpoint unless local integration is added. The Python-to-C++ converter supports only simple syntax, so real Python programs with classes, imports, comprehensions, lists, dictionaries, slices, exception handling, or NumPy calls will not convert directly. The feature extractor is regex-based, so it is fast and transparent but less semantically complete than full dependence analysis.

3.11 How to Use the Parallelizer

After installing LLVM/Clang and building the project bindings, a user can parallelize a C, C++, or supported Python file:

```
from accelera.src.utils.parallelizer import parallelizer

parallelizer.parallelize("examples/loop_example.py")
# Writes examples/parallelized_loop_example.c
```

A simple Python input that the converter can handle is:

```
total = 0
for i in range(0, 1000):
    total += i
print(total)
```

The converter emits C++; the parallelizer extracts the loop and can produce a result similar to:

```
#include <iostream>

int main() {
    int total = 0;
#pragma omp parallel for reduction(+ : total)
    for (int i = 0; i < 1000; i++) {
        total += i;
    }
    (std::cout << total << std::endl);
}
```

```

    return 0;
}

```

3.12 Hard-File Parallelizer Evaluation

The hard-file evaluation script `tools/evaluate_hard_parallelizer.py` was added to execute the parallelizer on the hard examples in `data/hard_test_files`. The script parallelizes each file in a temporary directory, compiles both the baseline and generated versions with OpenMP support, runs each executable three times, compares the outputs, and writes a JSON report to `hard_parallelizer_eval_results_generated.json`. For numeric outputs, the script uses tolerant floating-point comparison because OpenMP reductions can change the order of addition and therefore slightly change the last printed digits without changing the mathematical result.

Table 6: Hard-file parallelizer evaluation results.

File	Baseline time	Parallel time	Speedup
<code>hard_eval_000.cpp</code>	1786.28 ms	195.47 ms	9.138×
<code>hard_eval_001.cpp</code>	63.98 ms	31.47 ms	2.033×
<code>hard_eval_002.py</code>	42.66 ms	7.67 ms	5.563×

All three hard examples were parallelized, compiled, executed, and verified successfully. The generated and baseline outputs matched for every file. The summary report contained 11 extracted loops, 8 gold pragmas, and 8 generated pragmas. Exact pragma matching reached 8 out of 8 gold pragmas, and the average pragma similarity was 1.0. The average measured parallelization latency was 231.53 ms. Across the three verified files, the average speedup was 5.578×

3.13 Integration with Accelera Pipe

The parallelizer was also integrated with `accelera_pipe` so that custom preprocessing functions can be optimized automatically when their syntax is supported. The integration is implemented through the `Pipeline.preprocess` builder and the `parallelizer` utility. When a preprocessing object is detected as a custom Python function, the pipeline attempts to optimize it through `optimize_pymethod`. The optimization path is intentionally best-effort: if conversion, loop extraction, OpenMP annotation, or native compilation fails, the original Python function is preserved and the pipeline remains correct.

The integration path is:

1. Extract the Python source of the custom preprocessing function using `inspect.getsource`.
2. Convert the supported Python subset to C++ using `py2cpp_converter`.
3. Run the parallelizer in in-memory code mode `parallelize(code, file=False)`, which extracts loops without writing a generated `.c` file or loop JSON file.
4. Insert OpenMP pragmas such as `#pragma omp parallel for`.
5. Rewrite the generated function into a pybind11-compatible NumPy extension for two-dimensional `float64` arrays.
6. Compile the generated C++ code into a cached Python extension module under `.accelera.cache/compiled`.
7. Store the compiled function in the preprocess node, so the existing C++ graph execution can call it exactly like an ordinary Python callable.

The current compiler helper targets NumPy-style preprocessing kernels and uses `pybind11`, Python include paths, and OpenMP compiler flags. For the row-normalization example in `examples/parallel_accpipe.py`, the Python function:

```
def normalize_rows(X):
    for i in range(len(X)):
        s = 0
        for j in range(len(X[i])):
            s += X[i][j] * X[i][j]

        norm = s**0.5

        for j in range(len(X[i])):
            X[i][j] = X[i][j] / norm

    return X
```

is transformed into a `pybind11` function using mutable NumPy accessors and an OpenMP parallel loop. The generated kernel operates in-place on the input array and returns the same NumPy array object after normalization.

Table 7: Accelerate Pipe integration benchmark for automatic custom-function parallelization.

Mode	Samples	Reported pipeline time
Python custom preprocessing function	500,000	6.95 s
Compiled OpenMP preprocessing function	500,000	0.07 s

This experiment shows an approximate speedup of

$$\frac{6.95}{0.07} \approx 99.3\times$$

for the preprocessing-heavy example. Both runs used the same generated dataset shape, (500000, 25), the same first 500 samples as test data, and the same row-normalization logic. The result demonstrates that the parallelizer can serve not only as a source-to-source OpenMP tool, but also as a runtime acceleration backend for supported custom preprocessing functions in `accelera_pipe`.

4 AutoML Module

5 Deployment Module

6 Benchmark Platform Module

7 Module-Level View

The five report modules describe different layers of Accelera. The `accelera_pipe` module focuses on graph-based machine learning workflow construction and execution. The code parallelizer focuses on source loop analysis and OpenMP generation. The AutoML, deployment, and benchmark platform sections are reserved for later module descriptions.

Across the implemented sections, a common architectural pattern appears: Python exposes a convenient API and performs validation, while compiled or external services provide lower-level parsing, execution, classification, or infrastructure support.

8 Conclusion

The five-module structure illustrates the central engineering style of Accelera: expose simple Python interfaces while using compiled infrastructure, model services, and supporting platforms for performance-sensitive or operational work. In the current report draft, `accelera_pipe` is documented as a graph-oriented ML workflow system with branching, metric adaptation, path selection, executed graph reuse, and HTML reporting. The code parallelizer is documented as a static-analysis and OpenMP generation tool with a restricted Python converter, feature extraction, classifier-based loop selection, and local heuristic fallbacks. The remaining sections are intentionally left as space for future documentation.